

# DSC510: Introduction to Data Science and Analytics

## Lab 3: Data Visualization

Pavlos Antoniou

Office: B109, FST01



University of Cyprus  
Department of  
Computer Science



# Why data visualization?

---

- Raw data (files, databases) alone doesn't reveal clear associations or trends.
- Visualization translates data into a visual context, making it easier for humans to interpret.
- Helps uncover **patterns, trends, and outliers**.
- Visualization is a key part of Exploratory Data Analysis (EDA) to:
  - Discover patterns
  - Spot anomalies
  - Test hypotheses
  - Check assumptions
- Saves time and resources by quickly identifying which models are suitable for the data.

# Why visualization saves time & resources

---



- Regression vs Classification
  - A **scatter plot / pair plot** of x (input features) vs y (target variable) that shows a clear linear relationship → try Linear Regression.
  - If instead y (target variable) is categorical with clear clusters in the scatter plot → go for Classification models instead.
- Feature Relationships
  - A **correlation heatmap** that reveals that two variables are highly correlated → you can drop one and avoid wasting computation on redundant features.
- Data Distribution
  - A **histogram** that shows your target variable is highly skewed → instead of trying complex models right away, you can apply a transformation (log / box-cox / yeo-johnson) or choose models that handle skewness well.

# Data Visualization Libraries in Python

---

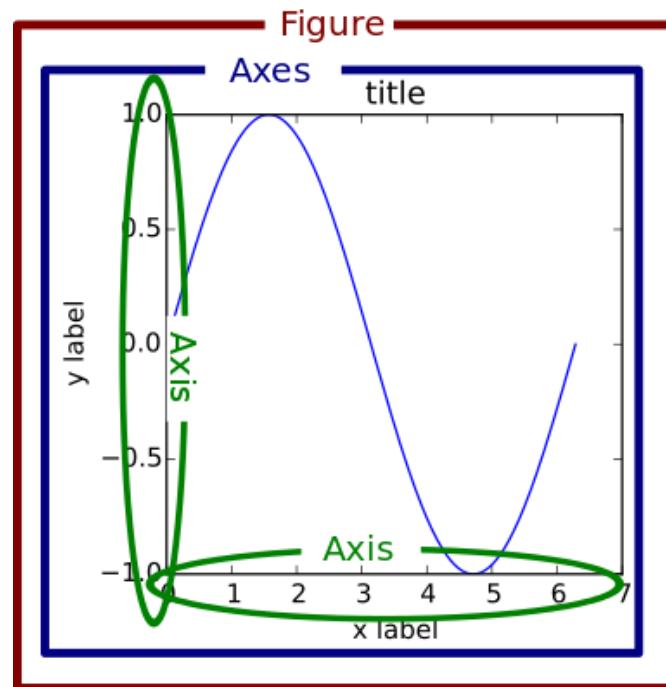


- Matplotlib
  - Most prominent plotting lib
  - import it to a notebook using: `import matplotlib.pyplot as plt`
- Seaborn
  - Visually appealing plots
  - Based on Matplotlib
  - import it to a notebook using: `import seaborn as sns`
- Plotly
  - Interactive charts and maps
  - Not pre-installed in Anaconda. Run: `conda install -c plotly plotly` on Anaconda prompt

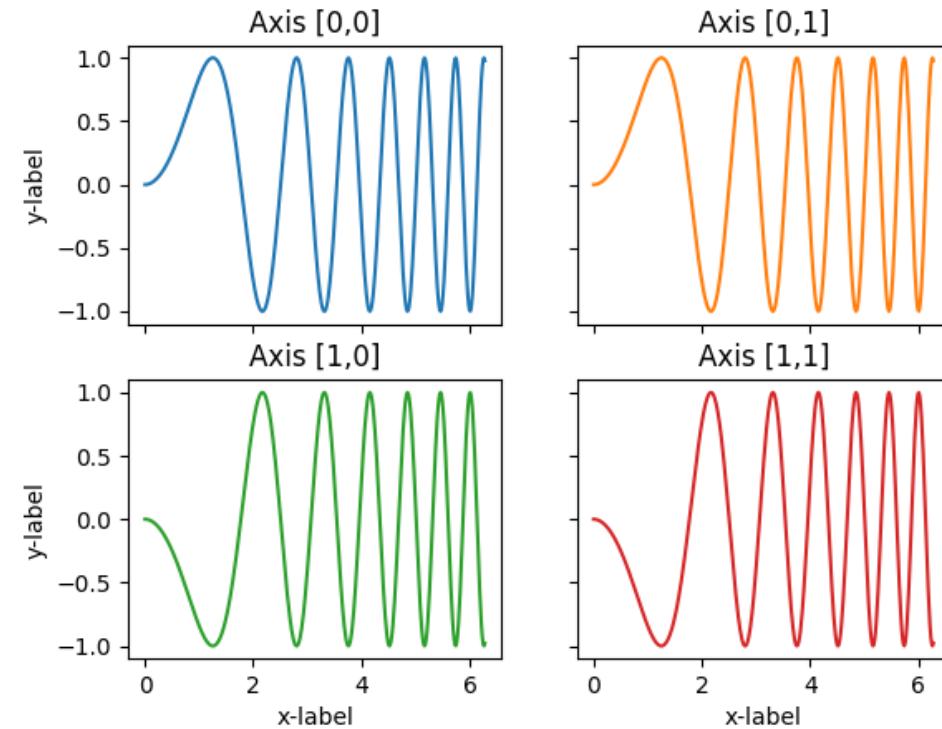


# Matplotlib Figure

- Matplotlib plots the data on **Figures** each of which can contain one or more **Axes**
- An **Axes** is attached to a Figure and contains a **region for plotting data and sets the coordinate system**



1 Figure with 1 Axes

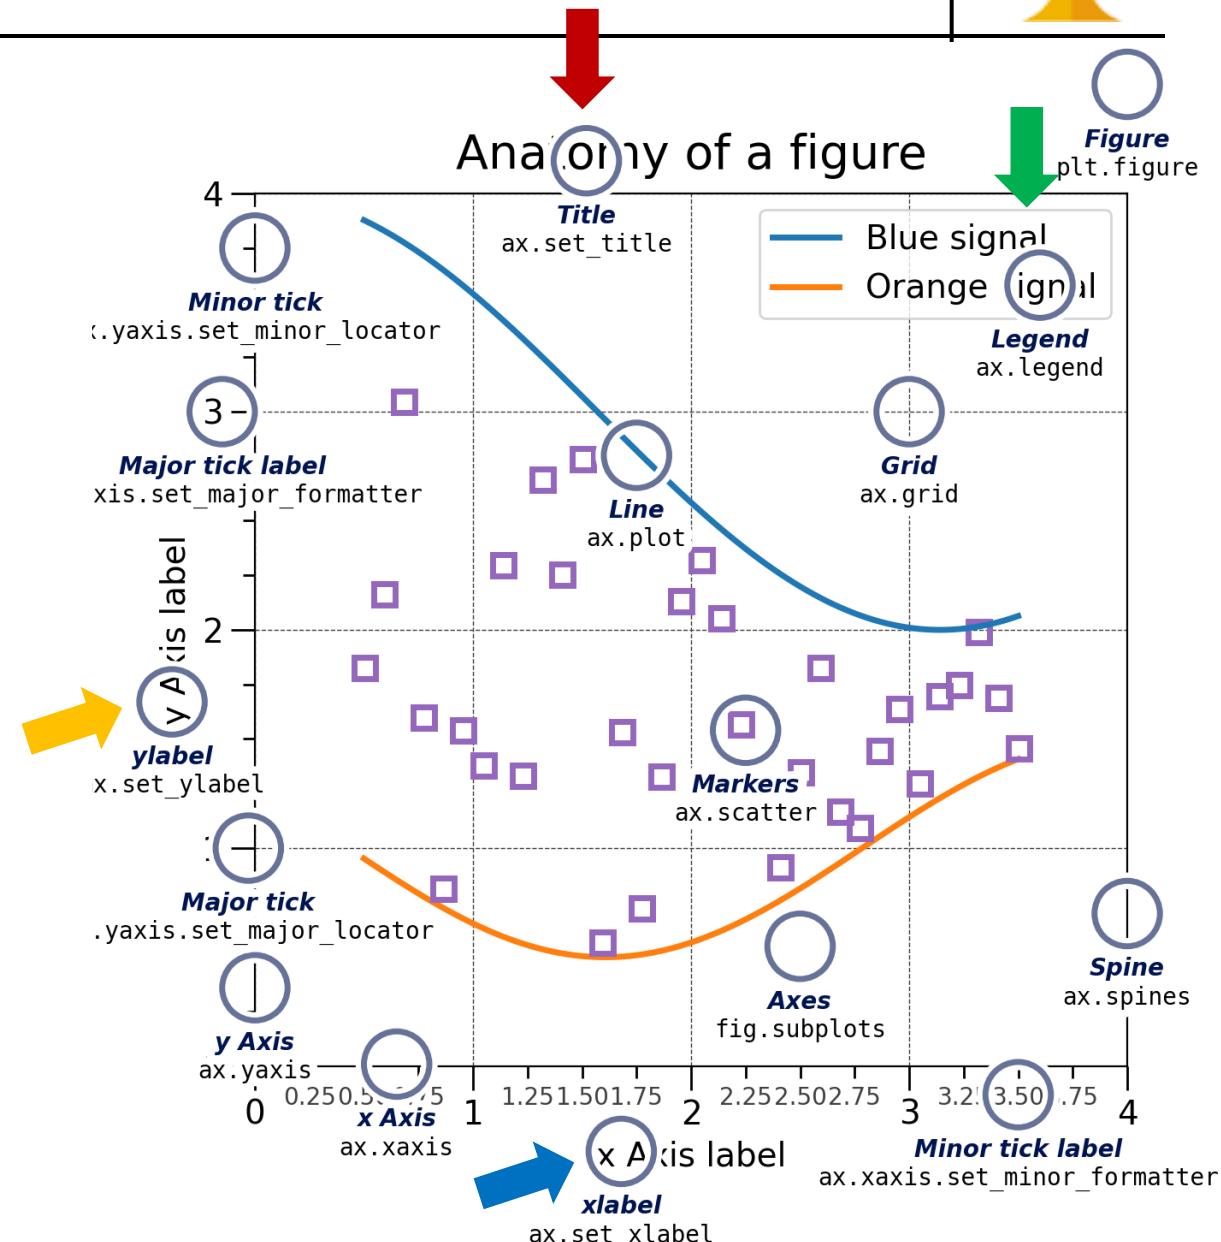


1 Figure with 4 (2x2) Axes



# Figure anatomy

- The Figure keeps track of all the child Axes, and of the group of 'special' objects (**titles**, figure **legends**, **xlabel**, **ylabel**, etc) belonging to each Axes





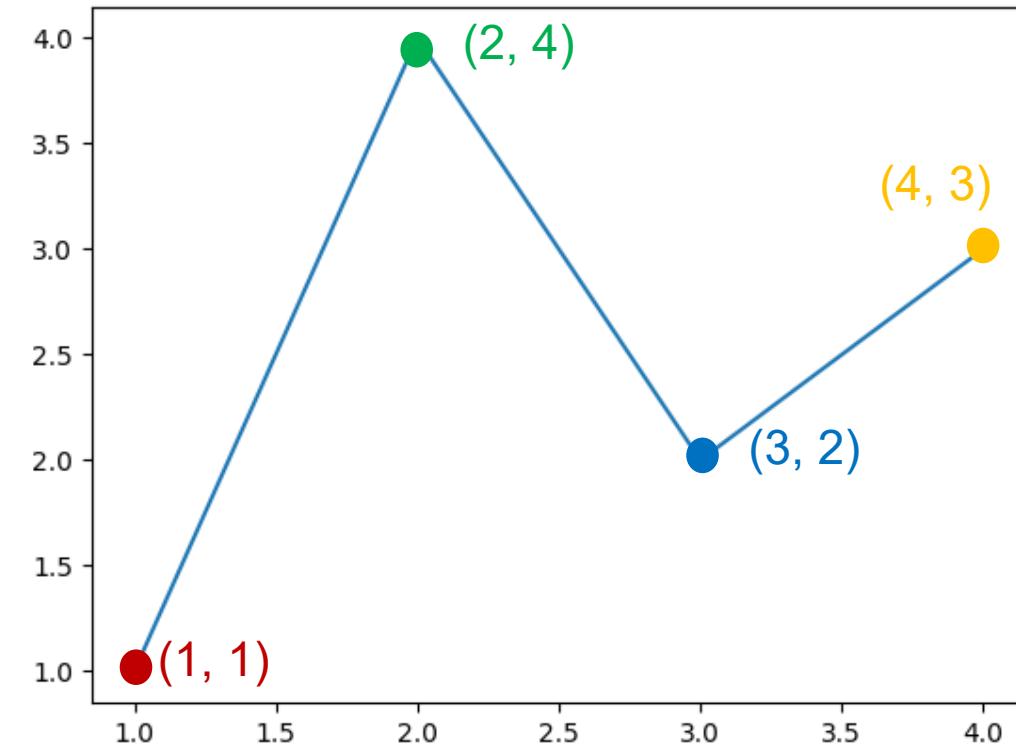
# Simple example

- The simplest way of creating a Figure is by using `plt.subplots` which returns an **Axes object**. We can then use `plot` function on the Axes object to draw some data:

```
# Create a figure containing a single axes
fig, ax = plt.subplots()
# Plot some data on the axes
ax.plot([1, 2, 3, 4], [1, 4, 2, 3])
```

- Alternatively, we can **create a figure with no axes object**. We can then use `plt.plot` to draw data

```
# Create a figure with no axes
plt.figure()
# Plot some data on the current figure
# an axes will be created automatically
plt.plot([1, 2, 3, 4], [1, 4, 2, 3])
```





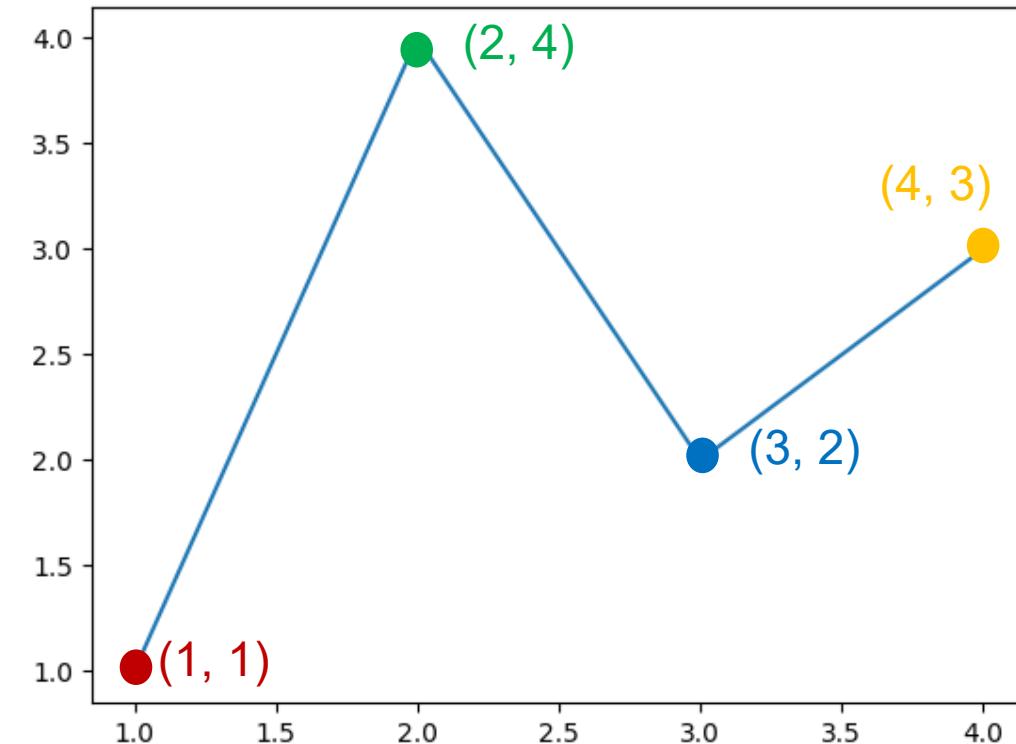
# Simple example

- The simplest way of creating a Figure is by using `plt.subplots` which returns an **Axes object**. We can then use `plot` function on the Axes object to draw some data:

```
# Create a figure containing a single axes
fig, ax = plt.subplots()
# Plot some data on the axes
ax.plot([1, 2, 4], [1, 4, 2, 3]);
```

- Alternatively, we can **create a figure with no axes object**. We can then use `plt.plot` to draw data

```
# Create a figure with no axes
plt.figure()
# Plot some data on the current figure
# an axes will be created automatically
plt.plot([1, 2, 3, 4], [1, 4, 2, 3]);
```





## `plt.subplots(nrows, ncols)`

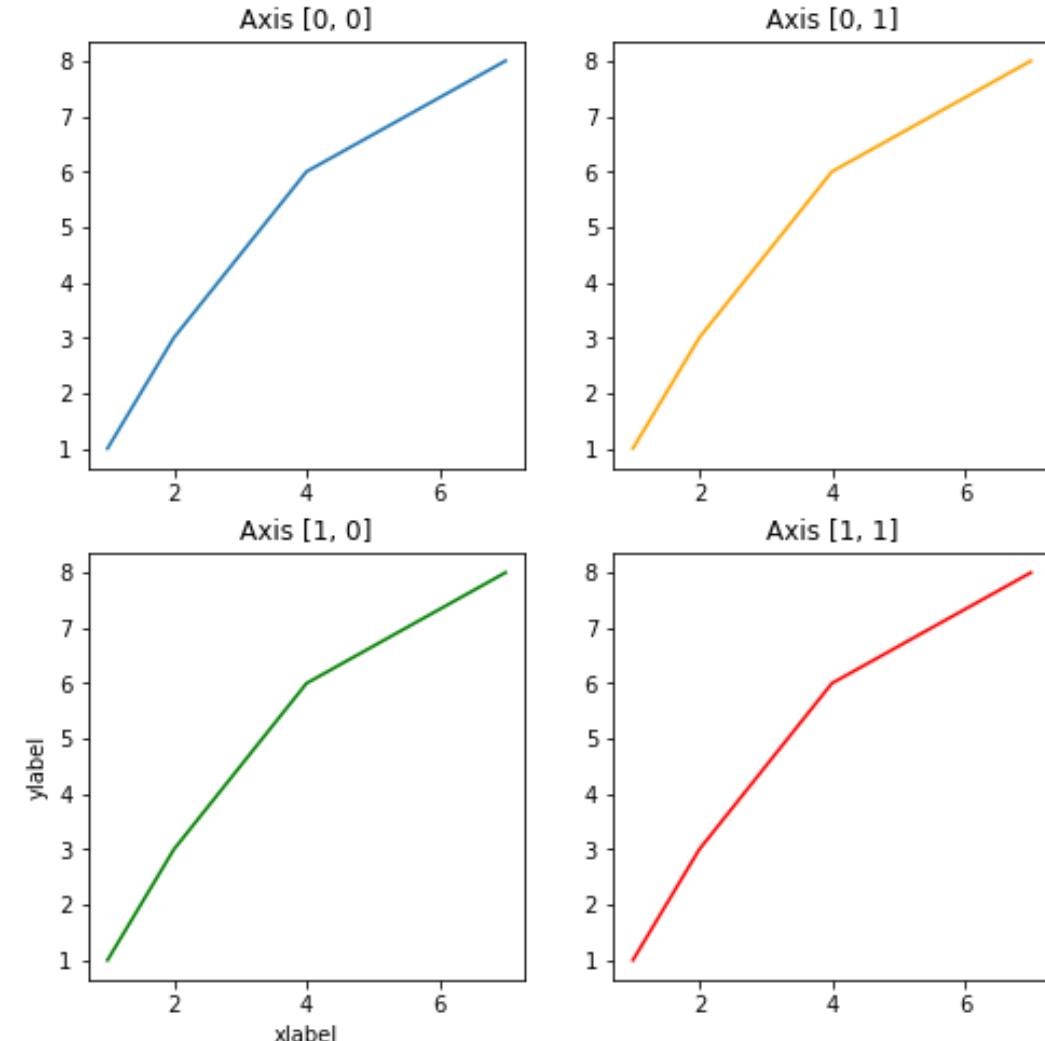
---

- Used when we want 2 or more axes plots in a single figure
  - One of the most useful features when we need to compare two or more plots hand to hand instead of having them separately
-



# Subplots – Option 1: Axes-level plotting

```
import matplotlib.pyplot as plt
import numpy as np
# Some example data to display
x = [1, 2, 4, 7] # x values
y = [1, 3, 6, 8] # y values
# Create a figure containing 2x2 axes
fig, axs = plt.subplots(2,2, figsize=(8,8))
# Plot some data on the axes
axs[0, 0].plot(x, y)
axs[0, 0].set_title('Axis [0, 0]')
axs[0, 1].plot(x, y, color='orange')
axs[0, 1].set_title('Axis [0, 1]')
axs[1, 0].plot(x, y, color='green')
axs[1, 0].set_title('Axis [1, 0]')
axs[1, 0].set_xlabel('xlabel')
axs[1, 0].set_ylabel('ylabel')
axs[1, 1].plot(x, y, color='red')
axs[1, 1].set_title('Axis [1, 1]')
```



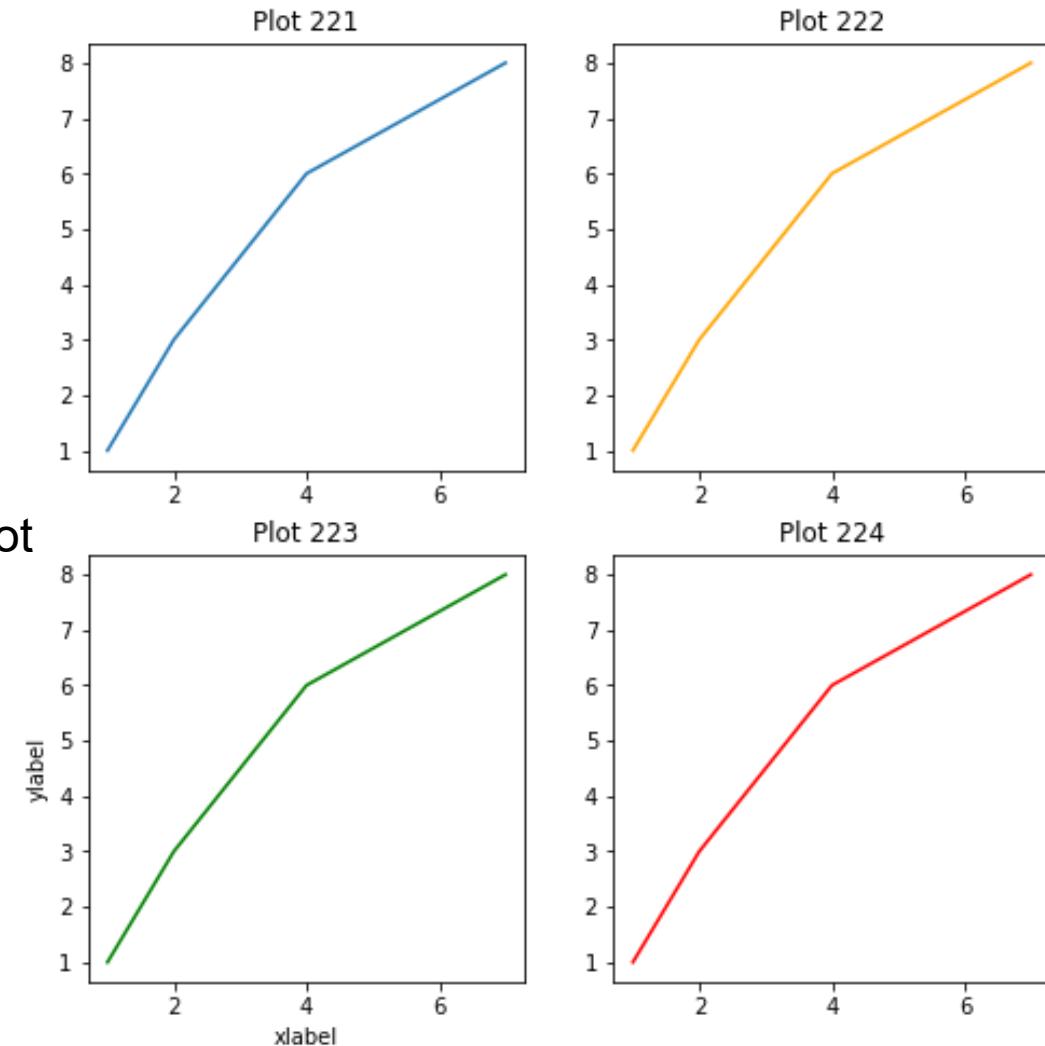
Specifies the width and height of the figure in unit inches.  
By default, the figure has the dimensions as (6.4, 4.8)

# Subplots – Option 2: Figure-level plotting



```
import matplotlib.pyplot as plt  
import numpy as np  
  
# Some example data to display  
x = [1, 2, 4, 7]  
y = [1, 3, 6, 8]  
  
# Create a figure with no axes  
plt.figure(figsize=(8, 8))  
  
# Plot some data  
plt.subplot(221)  
plt.plot(x, y)  
plt.title('Plot 221')  
  
plt.subplot(222)  
plt.plot(x, y, color='orange')  
plt.title('Plot 222')  
  
plt.subplot(223)  
plt.plot(x, y, color='green')  
plt.xlabel('xlabel')  
plt.ylabel('ylabel')  
plt.title('Plot 223')  
  
plt.subplot(224)  
plt.plot(x, y, color='red')  
plt.title('Plot 224')
```

3-digit integer where:  
1<sup>st</sup> num = No of rows  
2<sup>nd</sup> num = No of columns  
3<sup>rd</sup> num = index of that plot





# Different types of analysis

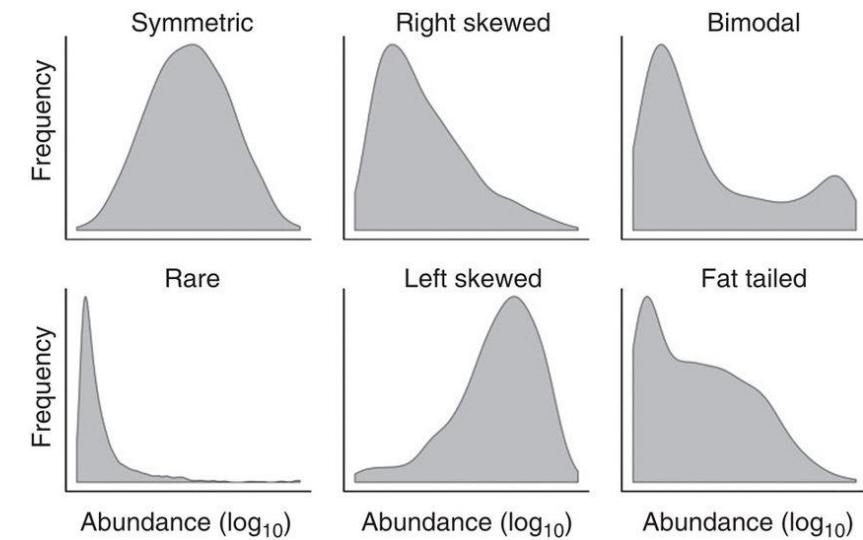
---

- There are different types of graphical analysis as mentioned below.
  - **Univariate**: In univariate analysis we are using a **single feature** (variable) to analyze almost of its properties (we create plots with one feature)
  - **Bivariate**: When we compare the data between exactly **2 features** then it's called bivariate analysis (we create plots involving 2 features)
  - **Multivariate**: Comparing **more than 2 features** is called as Multivariate analysis (we create plots with more than 2 features)
- Plot types discussed in the next slides will be marked as (U),(B) & (M) to represent them as Univariate, Bivariate and Multivariate plots correspondingly.



# Distribution plots (U/B)

- An early step in any effort to analyze or model data should be to understand how the variables are distributed
- Techniques for distribution visualization can provide quick answers to many important questions:
  - What range do the observations cover?
  - What is their central tendency?
  - Are they heavily skewed in one direction?
  - Is there evidence for bimodality?
  - Are there significant outliers?
  - Do the answers to these questions vary across subsets defined by other variables?





# Datasets for plotting

---

- In order to better show the usage of each plot we introduce two popular datasets:
  - Haberman dataset
  - Iris dataset



# Haberman (cancer survival) dataset

---

- Dataset contains 306 cases (rows or observations) from a study that was conducted between 1958 and 1970 at the University of Chicago's Billings Hospital on the survival of patients who had undergone surgery for breast cancer
- **Features** (= the **input variables**/columns of a dataset that are used by a predictive algorithm to make predictions):
  - Age of patient at time of operation (numerical)
  - Patient's year of operation (year - 1900, numerical)
  - Number of positive axillary nodes detected (numerical)
- **Target** (= **output variable** to predict or explain using features)
  - Survival status (categorical or class attribute)
    - 1 = the patient survived 5 years or longer
    - 2 = the patient died within 5 year



# Haberman dataset

---

```
df1 = pd.read_csv('haberman.csv', names=['age', 'op_year', 'axil_nodes', 'surv_status'])
```

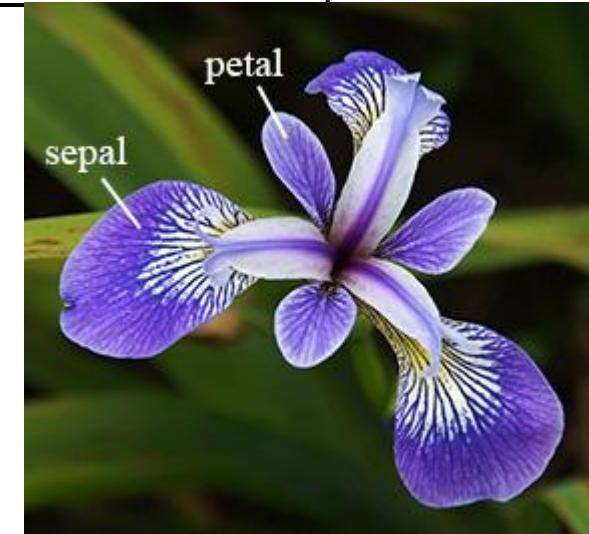
- DataFrame structure

	age	op_year	axil_nodes	surv_status
0	30	64	1	1
1	30	62	3	1
2	30	65	0	1
3	31	59	2	1
4	31	65	4	1
..	...	...	...	...
301	75	62	1	1
302	76	67	0	1
303	77	65	3	1
304	78	65	1	2
305	83	58	2	2



# Iris dataset

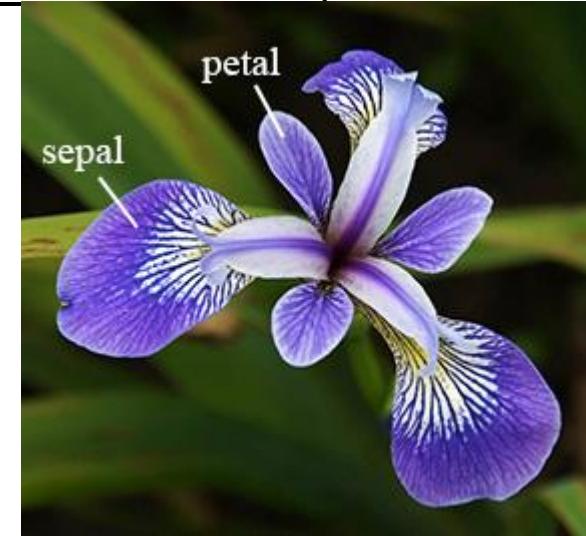
- Small dataset with 150 observations of iris flowers
  - each observation (row) has 4 columns of measurements (or variables or features) of the flowers (in centimeters)
  - target column (the 5<sup>th</sup> column) is the species (class) of the flower observed
    - all observed flowers belong to one of three species (setosa, versicolor, virginica)
- More info: [https://en.wikipedia.org/wiki/Iris\\_flower\\_data\\_set](https://en.wikipedia.org/wiki/Iris_flower_data_set)





# Dataset Overview

- Features:
  - sepal length in cm
  - sepal width in cm
  - petal length in cm
  - petal width in cm
- Target:
  - target column (class attribute)
    - Iris Setosa : 0
    - Iris Versicolour: 1
    - Iris Virginica: 2





# Iris dataset

---

```
df1 = pd.read_csv('iris.csv')
```

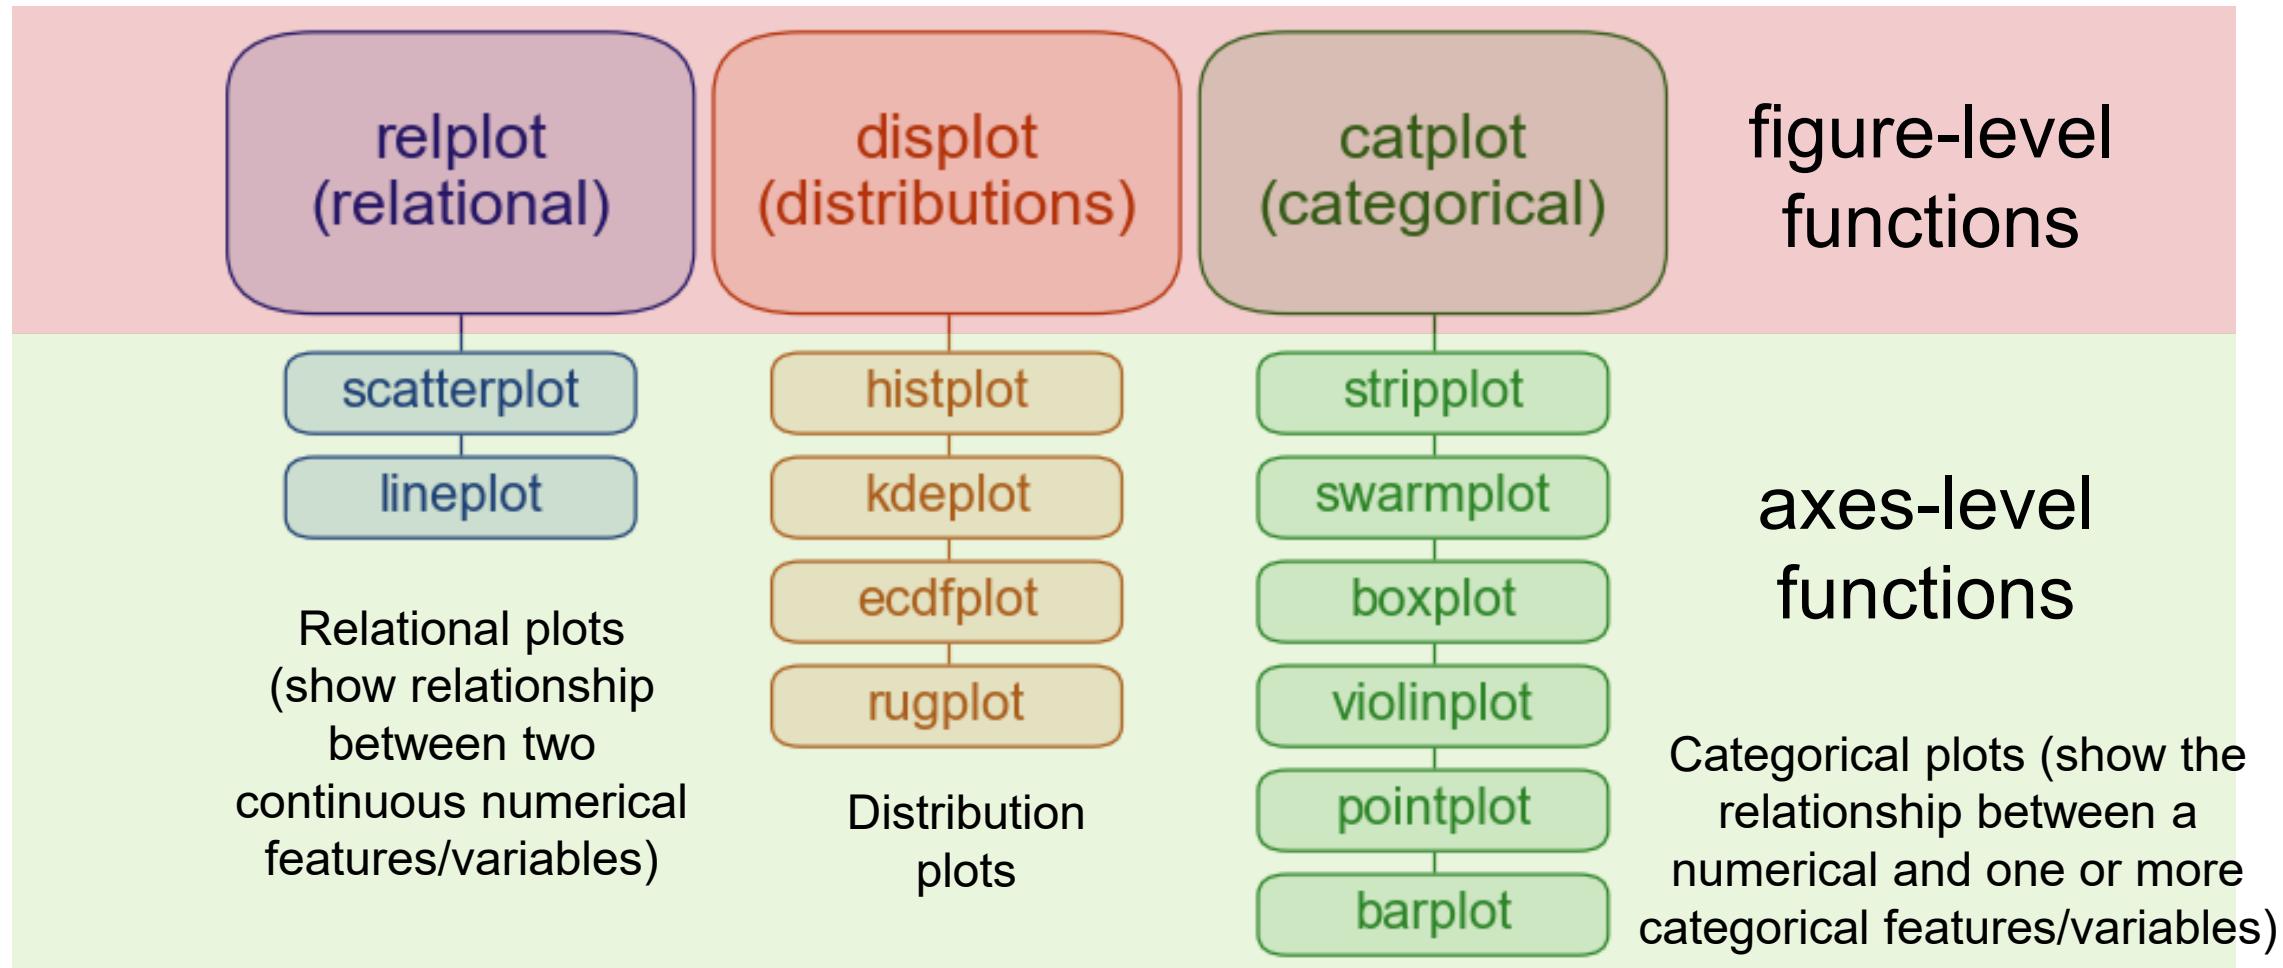
- DataFrame structure

	sepal_length_(cm)	sepal_width_(cm)	petal_length_(cm)	petal_width_(cm)	target
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0
..	...	...	...	...	...
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2
148	6.2	3.4	5.4	2.3	2
149	5.9	3.0	5.1	1.8	2



# Seaborn plotting functions

- Each seaborn module has a single figure-level function, which offers a unitary interface to its various axes-level functions





# Continuous (numerical) vs categorical data

---

- Continuous numerical variable/feature: contains data that can take on any value within a **defined range** and is often measured on a continuous scale, such as weight, height, or temperature
- Categorical variable/feature: contains data consisting of **discrete values** that fall into distinct categories or groups, such as gender, ethnicity, or product types. The values can be either strings or limited-range integer numbers
  - Example:
    - product type: electronics, food, furniture
    - product type: 0, 1, 2



# Distribution plots (U/B)

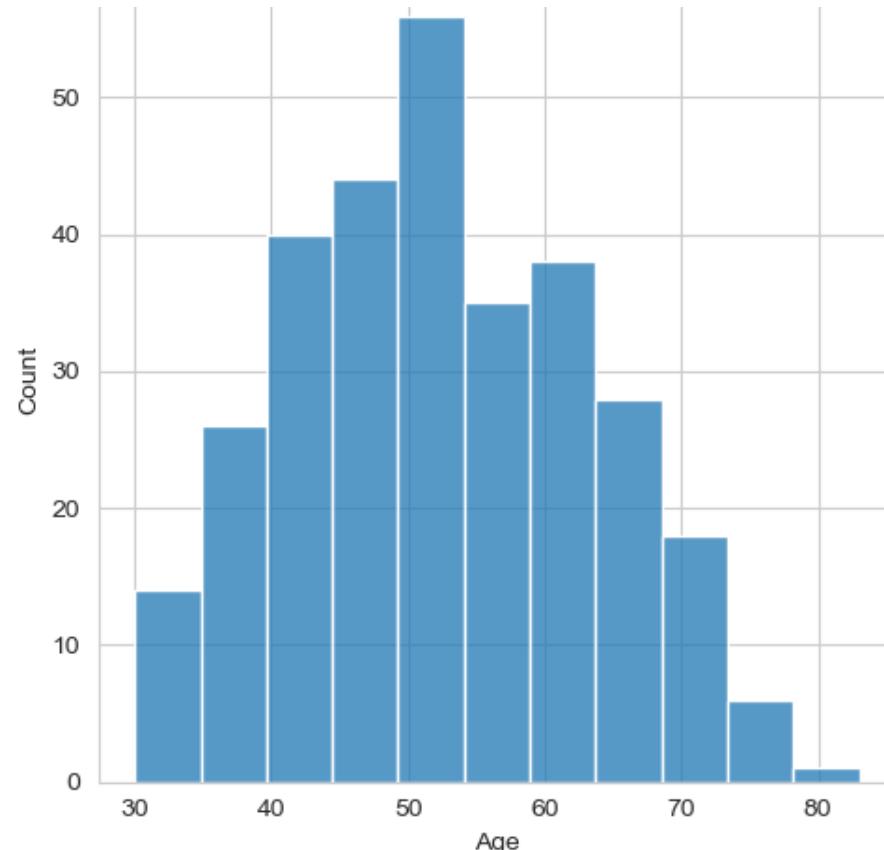
---

- Understand data distribution → tailor-made Machine Learning models to best fit our case study
  - Machine Learning models are designed to work best under some distribution assumption
    - ML models such as LDA, Gaussian Naive Bayes, Logistic Regression and Linear Regression perform better when all input variables (features) are normally distributed
  - Knowing with which distributions we are working with, can help us to identify which models are best to use or if we are in need of transforming data before applying any machine learning model
-



# Histogram plot (U) – displot() – Seaborn

- By default, displot() creates histogram (histplot() can be also used)
  - A histogram aims to approximate the underlying probability density function that generated the data by binning (grouping) and counting observations



```
plt.figure()  
sns.displot(data=df1, x='age')  
plt.xlabel('Age')
```

OR

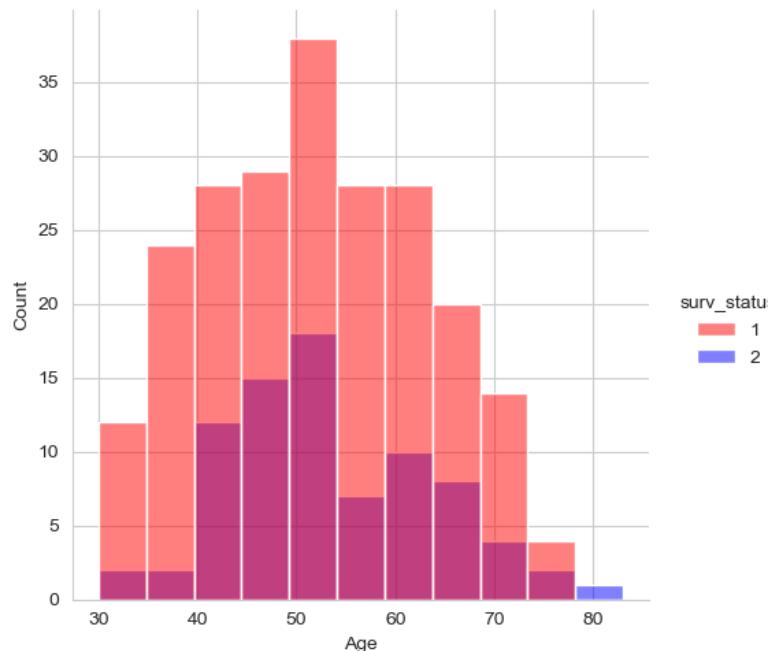
```
fig,ax = plt.subplots()  
sns.histplot(data=df1, x='age')  
ax.set_xlabel('Age')
```



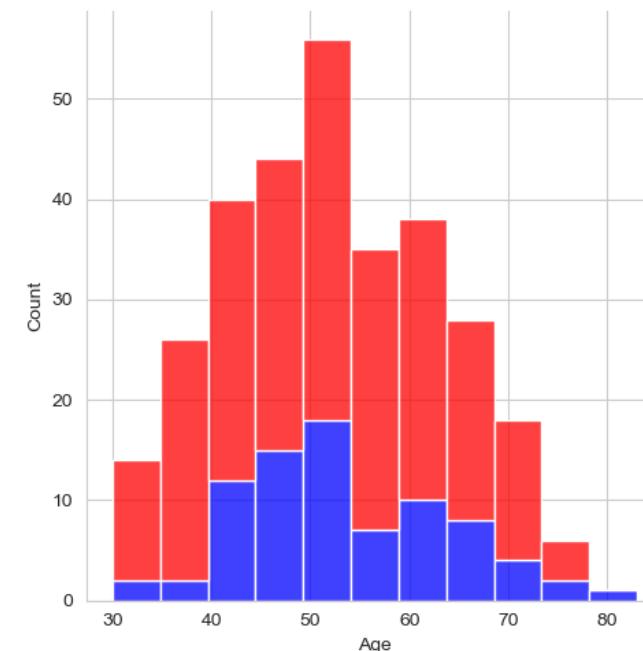
# Histogram plot (B) – displot() – Seaborn

- Once you understand the distribution of a variable, the next step is often to ask whether the behavior of that distribution differs across other variables in the dataset - *use of hue, usually with categorical var*

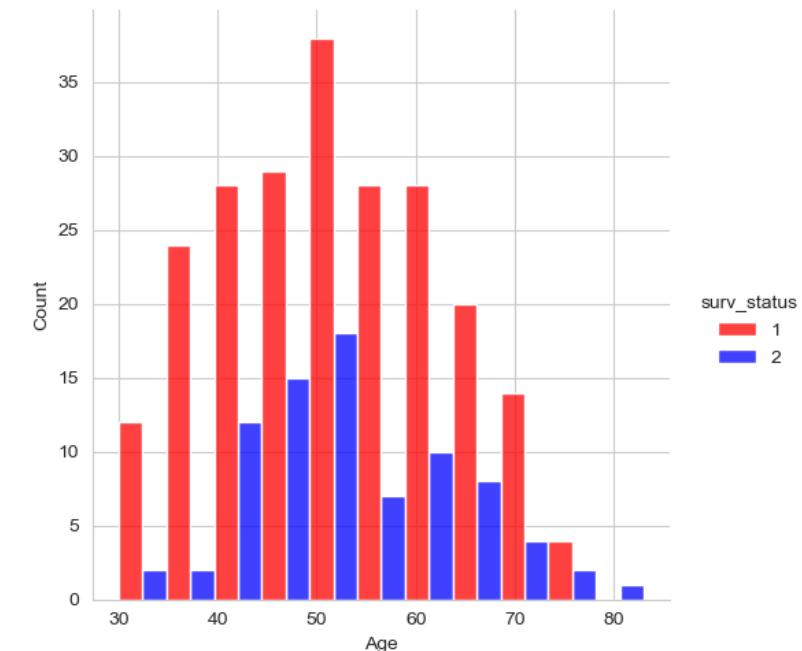
```
plt.figure()  
colors = ['red', 'blue']  
sns.displot(data=df1, x='age', hue='surv_status', palette=colors,  
plt.xlabel('Age')
```



multiple='stack'



multiple='dodge'



surv\_status  
— 1  
— 2



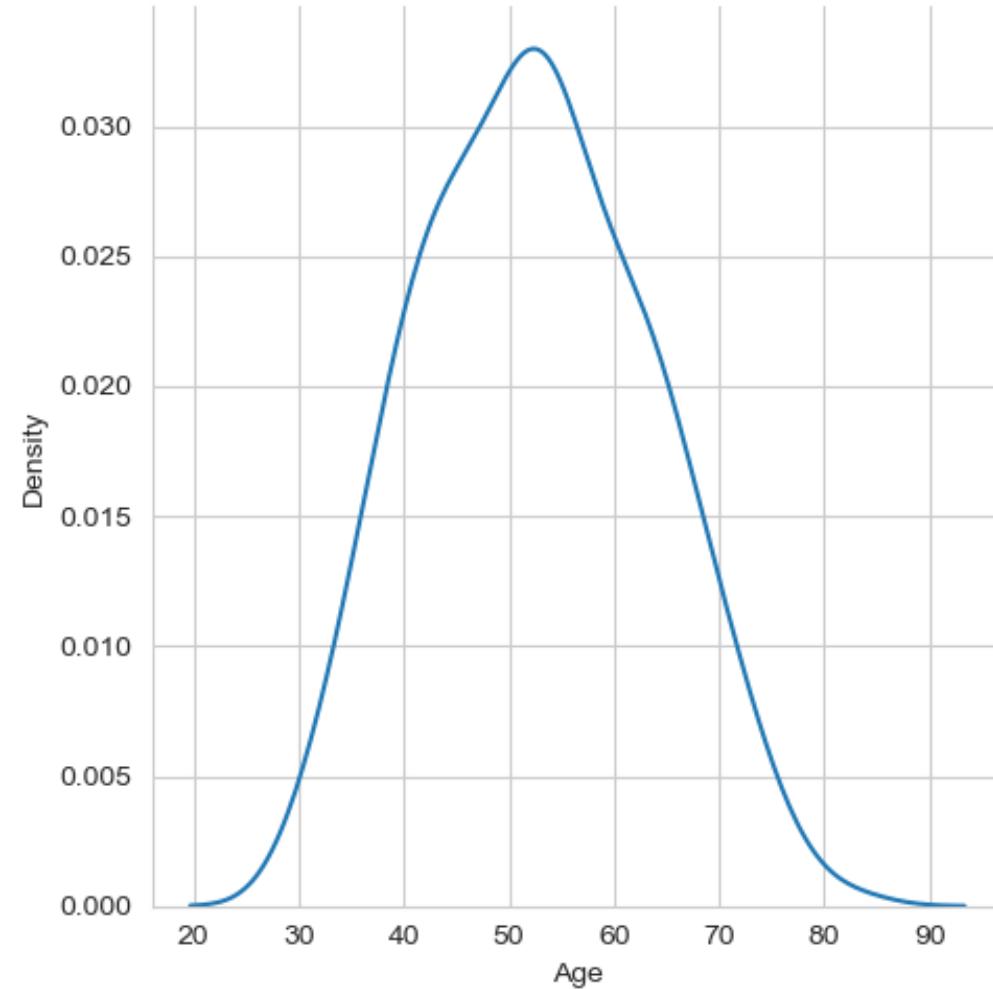
# Distribution plot (U) – displot() – Seaborn

- `displot()` with `kind='kde'`
  - same behavior as `kdeplot()`
  - rather than using discrete bins, a Kernel density estimation (KDE) plot smooths the observations with a Gaussian kernel, producing a continuous density estimate:

```
plt.figure()  
colors = ['red', 'blue']  
sns.displot(data=df1, x='age', kind='kde')  
plt.xlabel('Age')
```

OR

```
fig,ax = plt.subplots()  
sns.kdeplot(data=df1, x='age')  
ax.set_xlabel('Age')
```



# Percentiles



- *Percentile* a statistical measure that indicates the value below which a given percentage of observations in a dataset fall
  - Example: If you are the fourth tallest person in a group of 20
    - 80% of people are shorter than you:

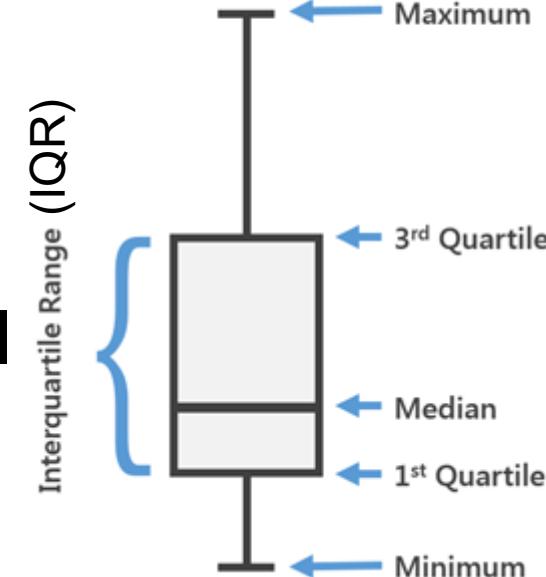


- That means you are at the **80th percentile**.
    - If your height is 1.85m then "1.85m" is the 80th percentile height in that group.
  - Q1 = 25th percentile = 1st quartile
  - Q2 = 50th percentile (*Median*) = 2nd quartile
  - Q3 = 75th percentile = 3rd quartile
  - Q4 = 100th percentile = 4th quartile
  - IQR (inter quartile range) = Q3 – Q1 (middle half of the observations)



# Box plots (or box-and whisker plots) (U)

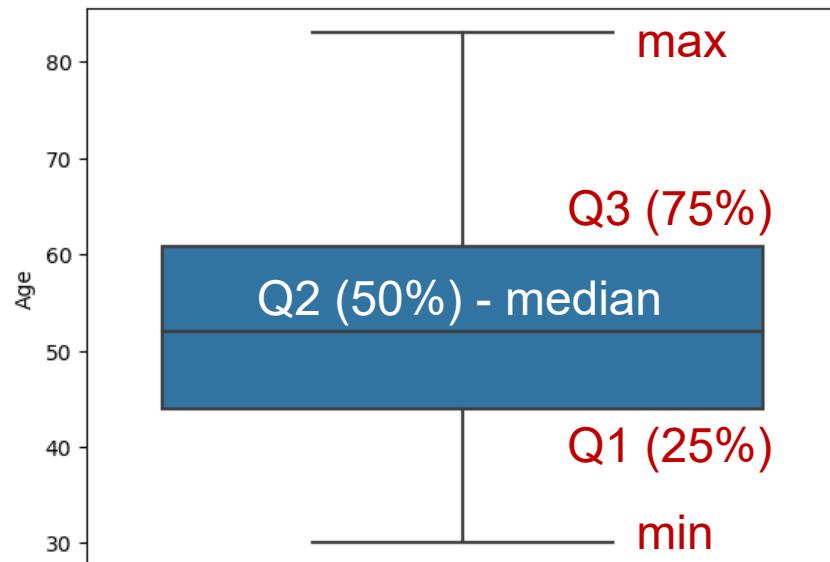
- Used to visualize the **statistical details** of a feature
- Box plots also give a clear overview of **outlier points**
- Box visualizes the interquartile range (IQR)
  - contains the middle **half** of the data set
- Straight lines at the statistical maximum and statistical minimum are called as whiskers
  - Statistical Maximum:  $Q3 + 1.5 * IQR$
  - Statistical Minimum:  $Q1 - 1.5 * IQR$
- Points outside of whiskers can be inferred as outliers
- The box plot gives us a representation of 25<sup>th</sup> percentile (or 1<sup>st</sup> quartile), 50<sup>th</sup> percentile (or 2<sup>nd</sup> quartile or median), 75<sup>th</sup> percentile (or 3<sup>rd</sup> quartile)





# Box plots (U) – Seaborn

- boxplot() function is available in the seaborn library
  - data parameter: dataset for plotting
  - x, y, hue parameters: names of features (variables) in data
- Box plots offer univariate analysis when we are exploring one feature (variable), however, multivariate analysis can be performed (see the extended set of slides)



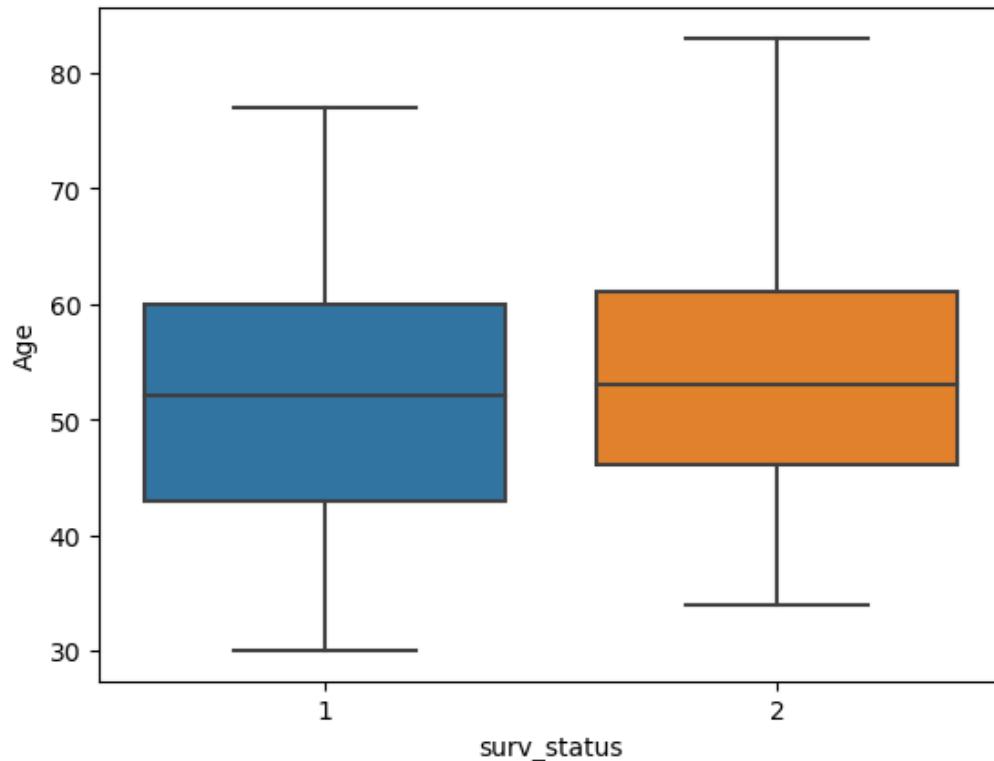
```
fig, ax = plt.subplots()  
sns.boxplot(data=df1, y='age')  
ax.set_ylabel('Age')
```

```
print(df1['age'].describe())  
count      306.000000  
mean       52.457516  
std        10.803452  
min        30.000000  
25%        44.000000  
50%        52.000000  
75%        60.750000  
max        83.000000
```



# Box plots (B) – Seaborn

- Assign a variable to x-axis to examine the statistical details for a combination of two variables



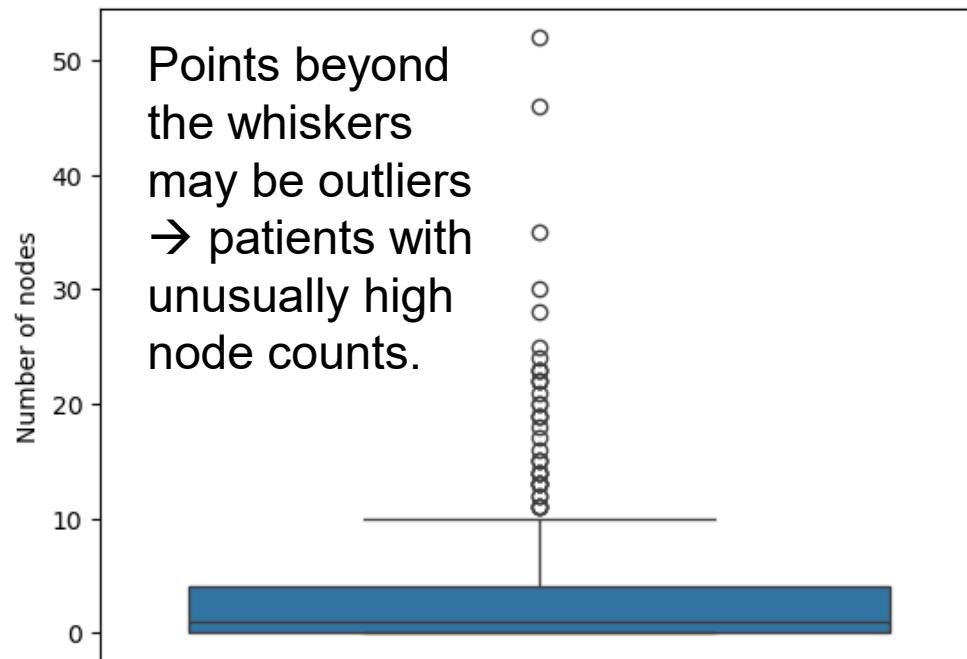
```
fig, ax = plt.subplots()  
sns.boxplot(data=df1, y='age', x='surv_status')  
ax.set_ylabel('Age')
```



# Box plots for outlier detection

```
fig, ax = plt.subplots()  
sns.boxplot(data=df1, y='axil_nodes')  
ax.set_ylabel('Number of nodes')
```

- Outliers are points that lie outside the whiskers.
- Box plots allow quick spotting of unusual values that **may**:
  - Indicate data entry errors
  - Represent rare but valid cases
  - Skew averages and affect models
- Important to detect them before training





# How to handle outliers?

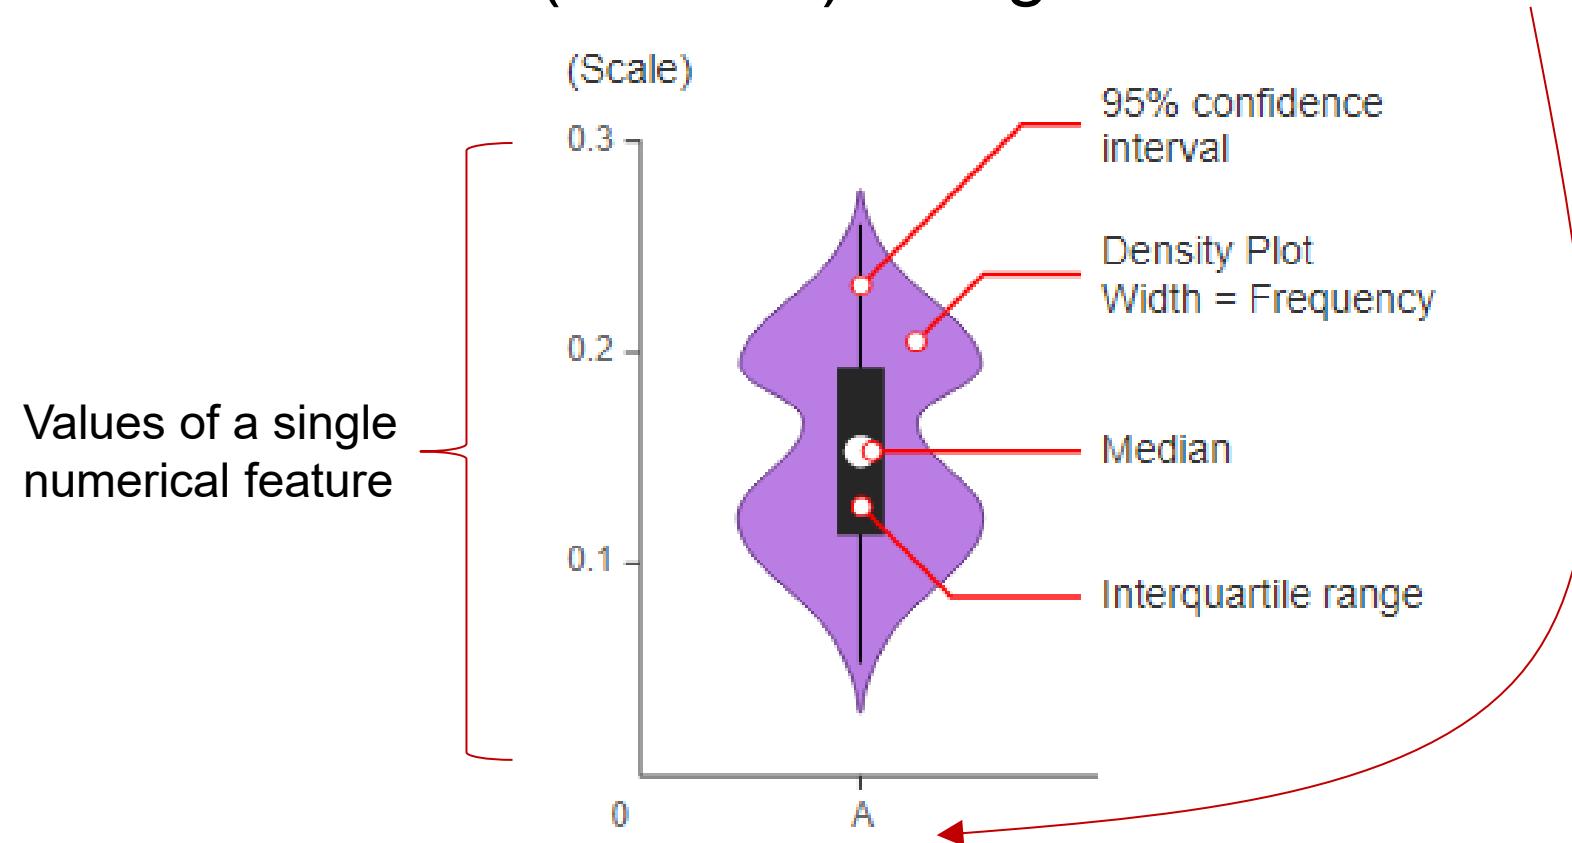
---

- Investigate First (Don't drop blindly)
  - Outliers may represent data entry errors (e.g. age > 120). → Fix or remove.
  - They may be genuine rare cases (e.g., patients with 50+ nodes in Haberman dataset). → Keep if medically meaningful.
- Options to Handle Outliers:
  - Drop: Safe if they are clearly errors or irrelevant to the analysis.
  - Transform: Apply log/square root transformations to reduce their influence.
  - Use robust models: Some algorithms (e.g., tree-based methods, median-based statistics) are less sensitive to outliers.
- Impact on Models
  - Dropping outliers may improve performance if they are noise.
  - But if they represent a real subgroup (e.g., patients with very severe cases), removing them may cause the model to miss important patterns.



# Violin plots (U/B/M) – Seaborn

- Shows the same summary statistics (median, IQR) as box plots
- Also show shape/distribution of a single numerical feature across several levels of one (or more) categorical variables



# Scatter plot (B) – Matplotlib / Seaborn / Plotly

---



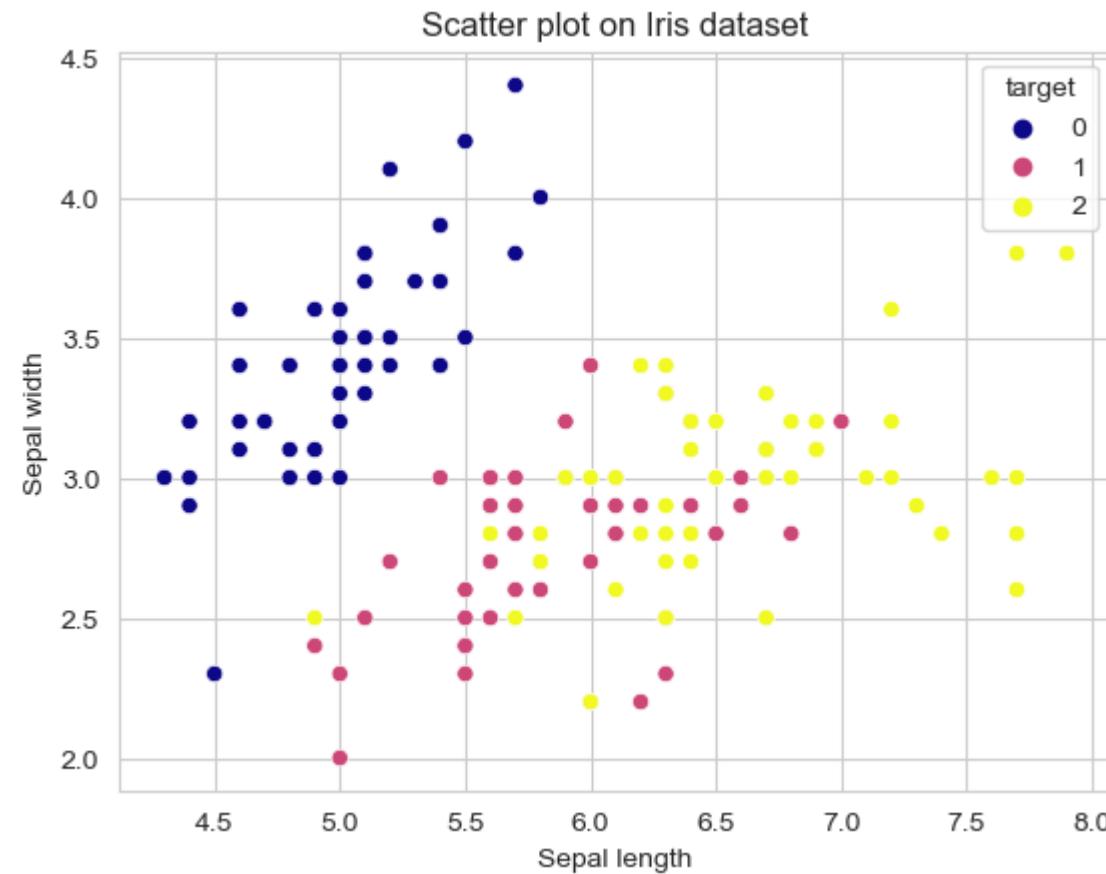
- A scatter plot shows where each data point (observation) lies with respect to any 2 or 3 features (dimensions).
  - Scatter plots can be drawn in 2D or 3D.
  - 2D scatter plot is primarily used to **find patterns/clusters** and **separability of the data** (this helps in discovering **important features**).
-

# Scatter plot (B) – Seaborn



```
fig, ax = plt.subplots()  
sns.scatterplot(data=df2, x = "sepal_width_(cm)", y = "sepal_length_(cm)", hue =  
"target", palette="plasma")
```

Parameter **palette** decides  
the colormap.



Parameter **hue** decides the  
color of each datapoints.



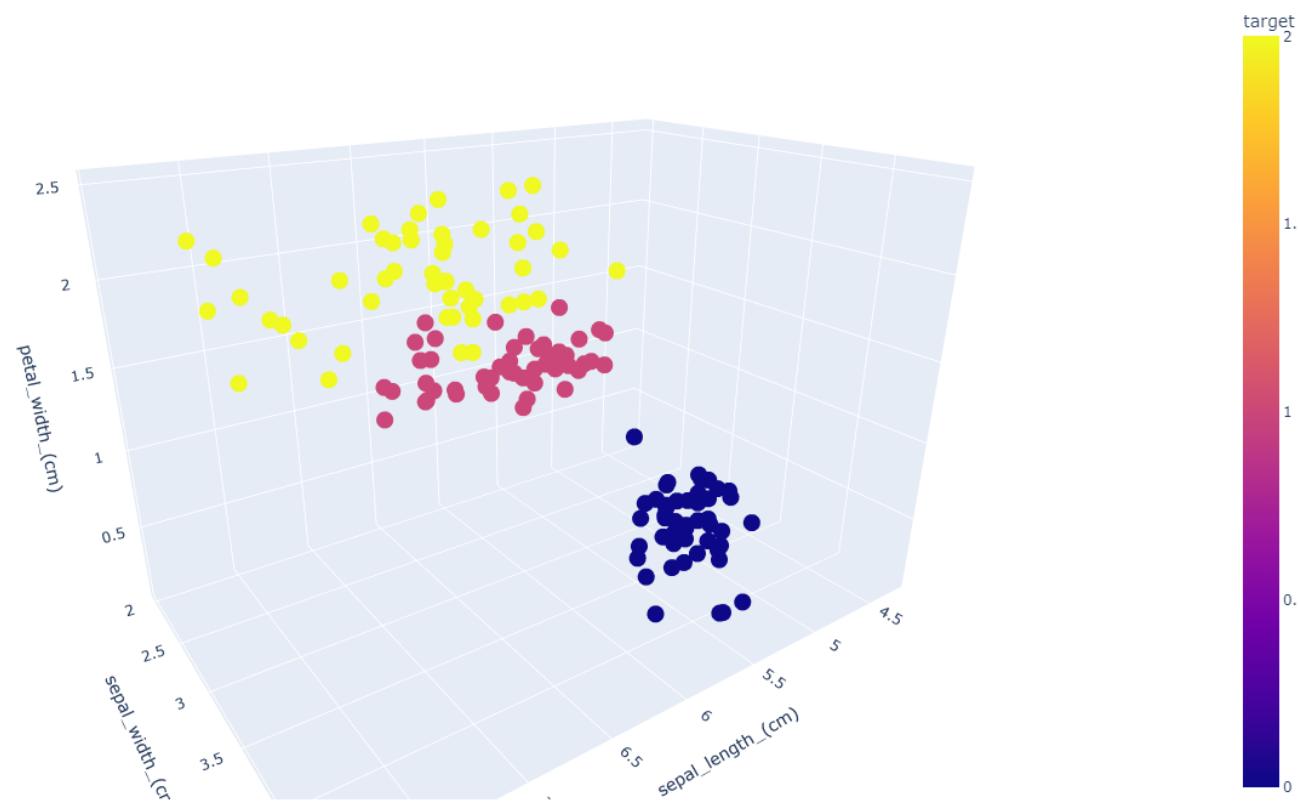
# 3D Scatter plot (M) – Plotly

- 3D scatter plot with Plotly Express

```
import plotly.express as px
fig = px.scatter_3d(df2, x='sepal_length_(cm)', y='sepal_width_(cm)',
z='petal_width_(cm)', color='target')
fig.show()
```

Parameter `color` decides the color of each datapoints.  
Same as `c` in matplotlib and `hue` in seaborn

Not pre-installed in Anaconda:  
conda install -c plotly plotly

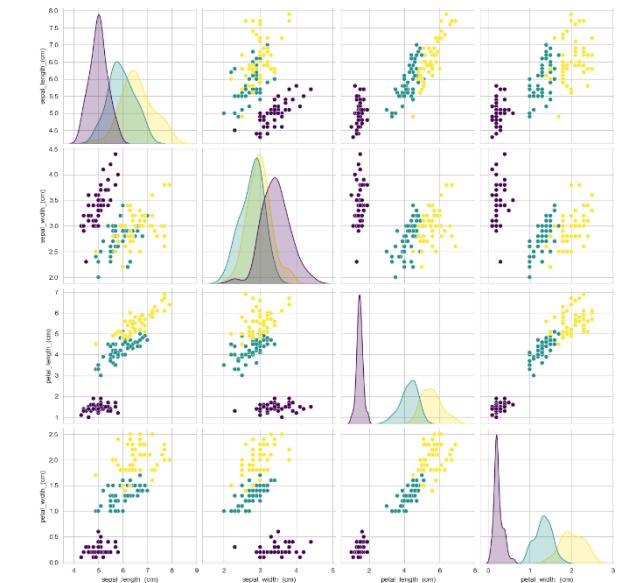




# Pair plots (M) – Seaborn

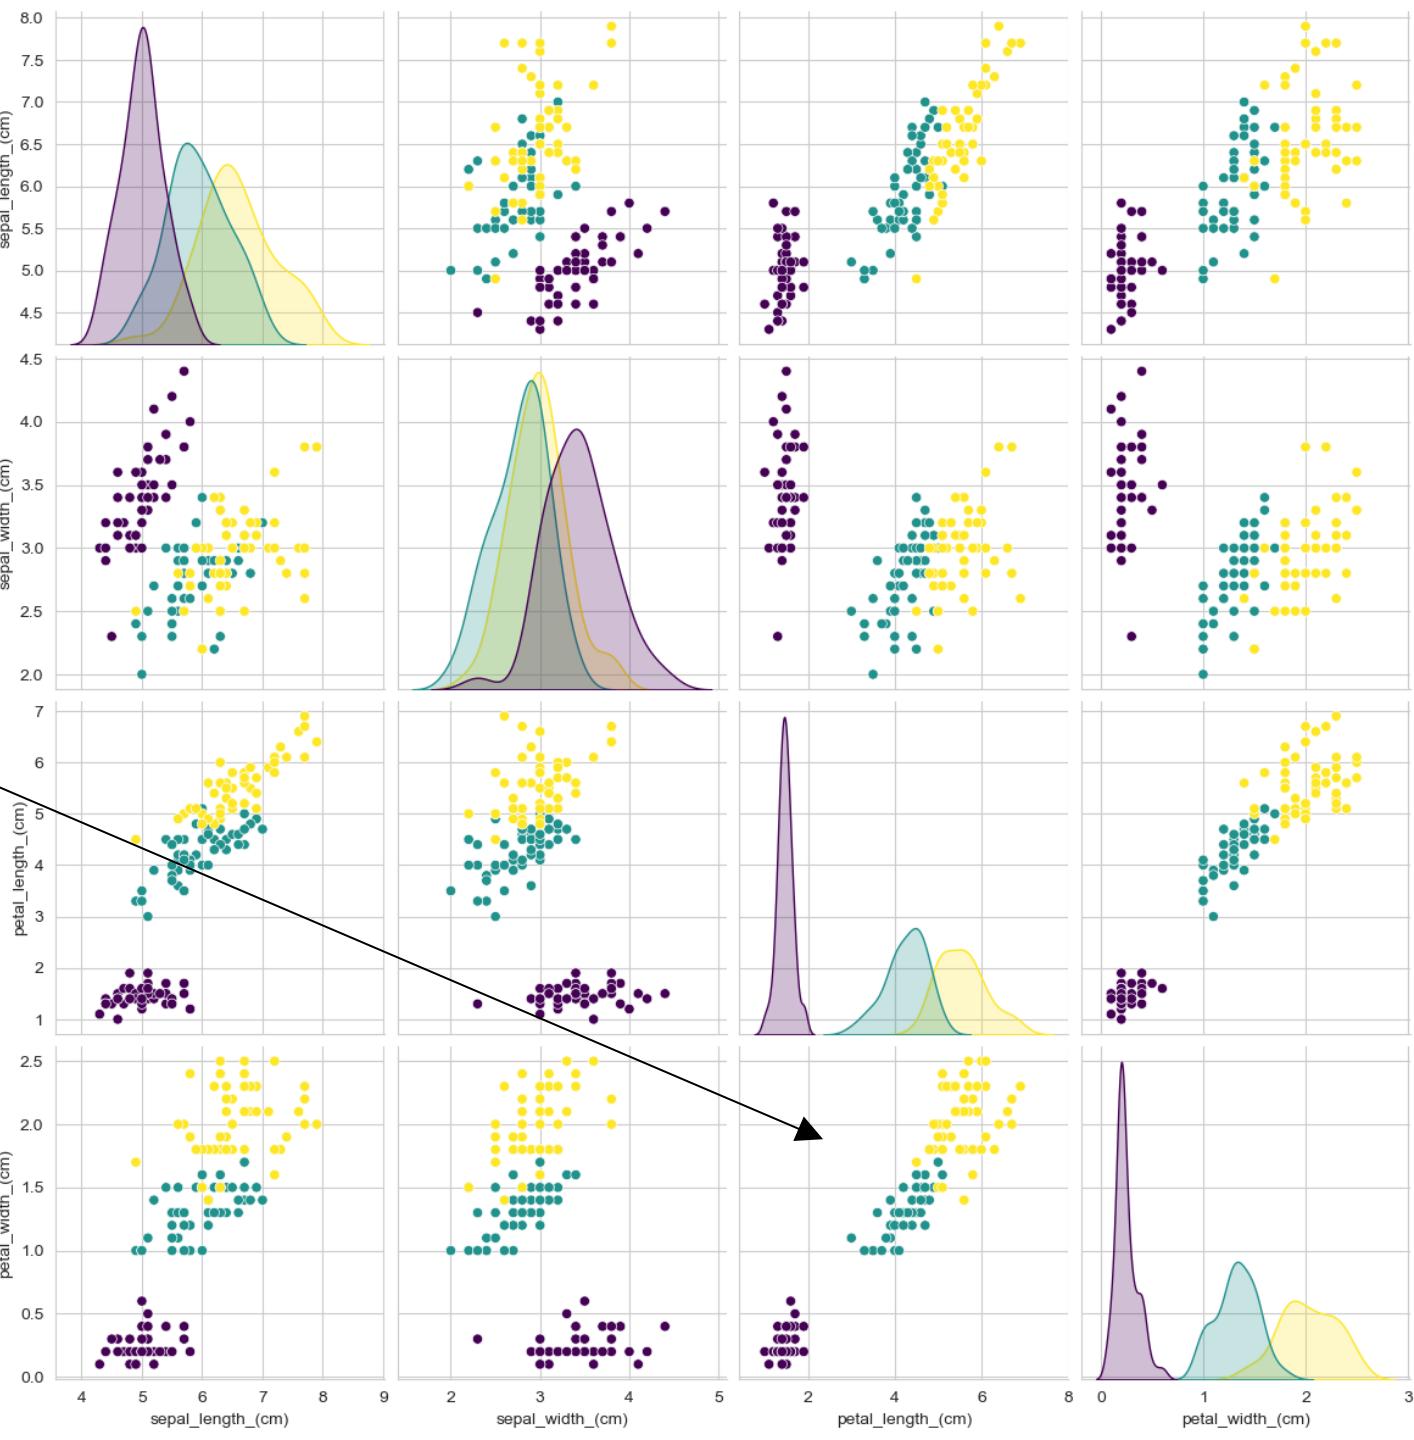
- Pair plot from seaborn: for scatter plots 4D and over
- For n features, pair plot will create a  $n \times n$  figure where the diagonal plots will be univariate distribution plot of the feature corresponding to that row and rest of the plots are the combination of features from each row in y axis and feature from each column in x axis.

```
plt.figure()  
sns.pairplot(data=df2, hue='target', palette='viridis')
```



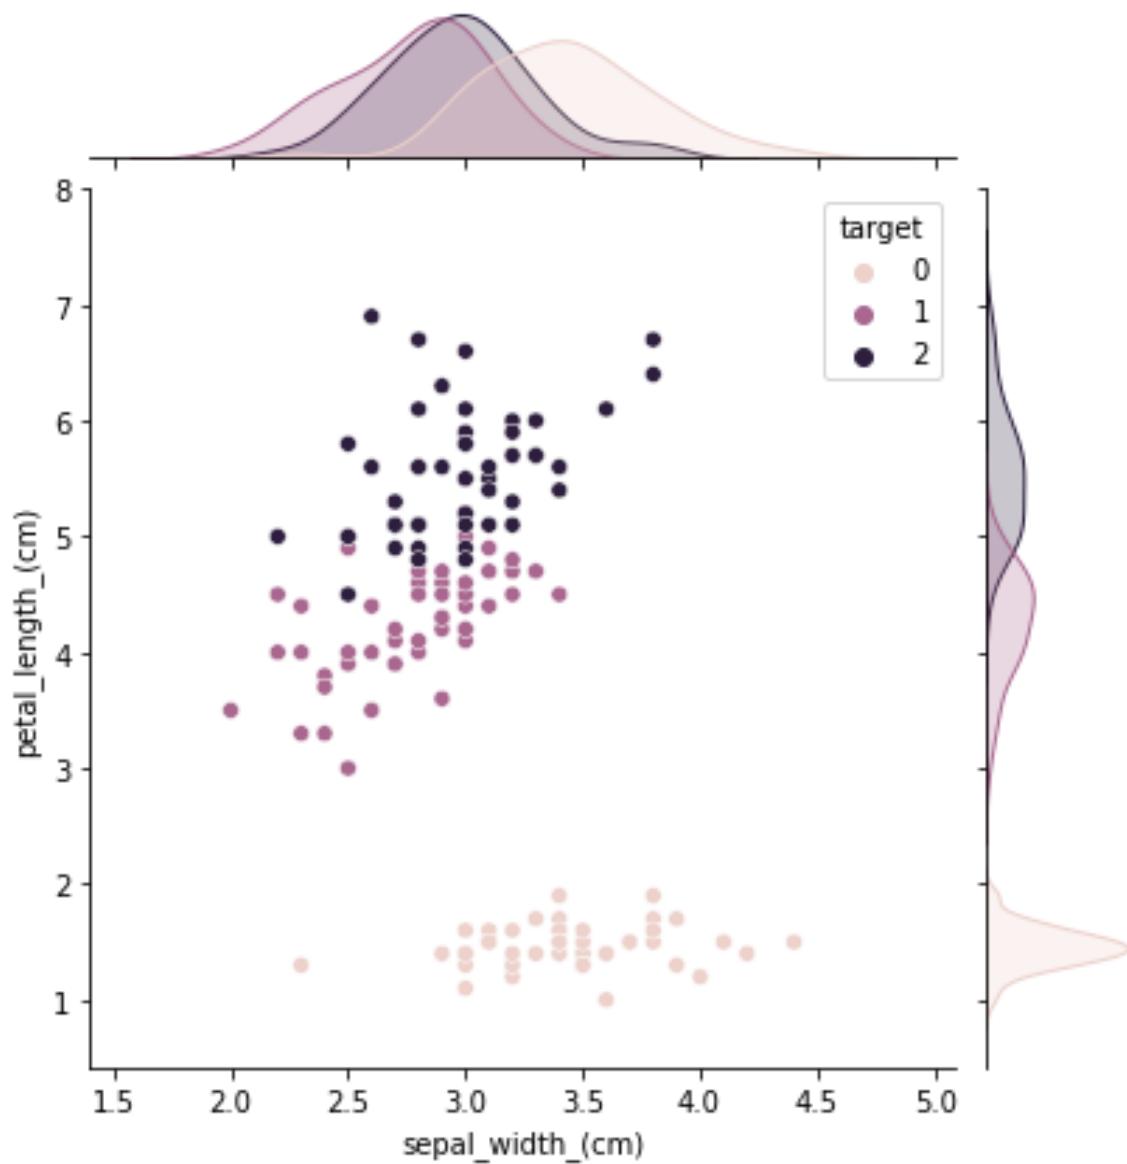
# Pair plots (M) – Seabc

- We can observe which 2 features can well explain/separate the data
  - It seems petal length and petal width are the 2 features which can separate the data very well  
=> important features
- Becomes complex in high number of features → dimensionality reduction techniques can be used (see Lab 5)



# Joint plot (U/B) – Seaborn

- Seaborn provides jointplot()
- Plot in the middle involves bivariate analysis whereas on the top and right side provides univariate plots of both variables
  - By default, jointplot() represents the bivariate distribution using scatterplot() & marginal distributions using histplot()
- Assigning a hue variable will add conditional colors to the scatterplot and draw separate density curves on the marginal axes



```
sns.jointplot(data=df2, x='sepal_width_(cm)', y='petal_length_(cm)', hue='target')
```

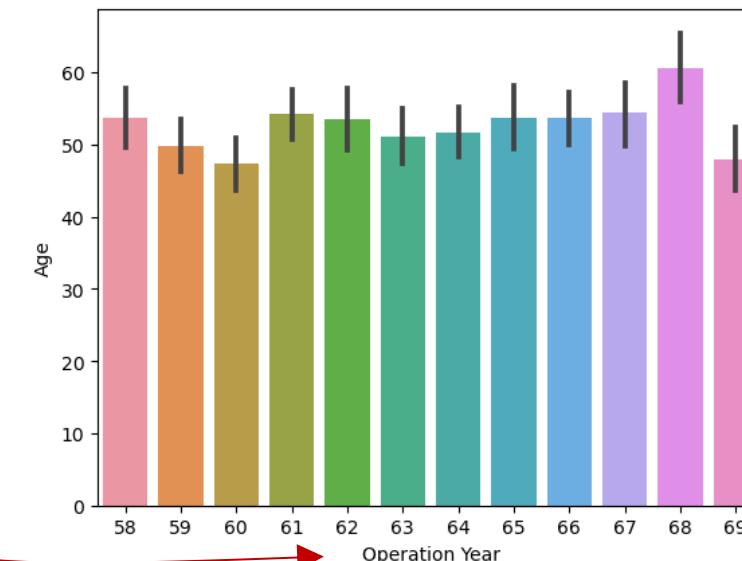


## Bar plot (B) – Seaborn

- Presents the **values** of a **categorical** (or discrete numerical) **variable** with **rectangular bars** with **heights proportional to a statistical measure** (mean, sum, median) of a **numerical** variable
- The height of the bar represents the value of the statistical measure
  - numeric value is estimated by aggregating across multiple observations of the y (numeric) variable at the same x (categorical) level – default is mean
  - indication of uncertainty (variation) around that value provided using error bars
    - Can be disabled using parameter ci=None

	age	op_year	axil_nodes	surv_status
0	30	64	1	1
1	30	62	3	1
2	30	65	0	1
3	31	59	2	1
4	31	65	4	1

category  
y='age', x='op\_year'  
Aggregate (group) by op\_year and estimate the mean value of aggregated ages

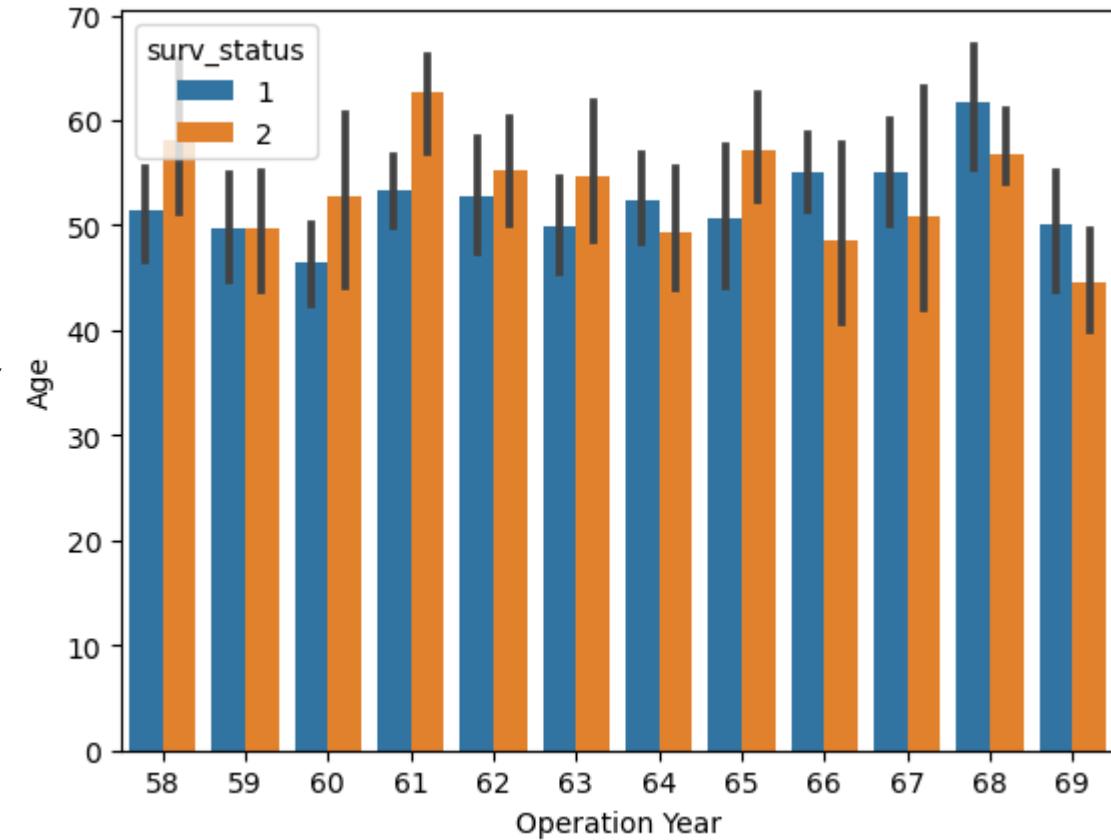




# Bar plot (M) – Seaborn

- Assigning a variable to hue will draw a separate bar for each of its unique values and distinguish them by a different color

```
fig, ax = plt.subplots()  
sns.barplot(data=df1, y='age', x='op_year',  
hue='surv_status')  
ax.set_xlabel('Operation Year')  
ax.set_ylabel('Age')
```



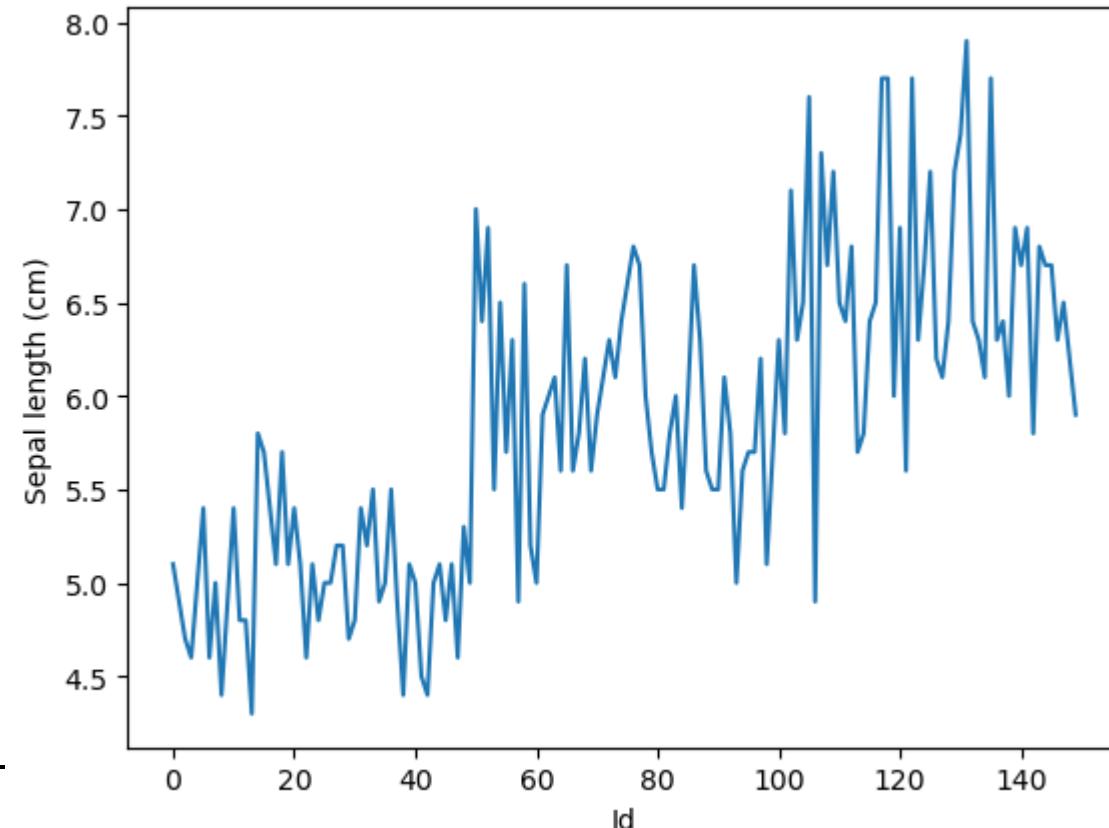


# Line plot (B) – Matplotlib & Seaborn

- A graph that displays data as points above a number line
- For variables (features) that can be ordered across another variable
  - Useful for timeseries data, where x-axis is a time-dependent variable (i.e. date)
- plt.plot() in Matplotlib
- sns.lineplot() in Seaborn

```
fig, ax = plt.subplots()  
ax.plot(df2.index, df2['sepal_length_(cm)'])  
ax.set_xlabel('Id')  
ax.set_ylabel('Sepal length (cm)')
```

```
fig, ax = plt.subplots()  
sns.lineplot(data=df2, x=df2.index,  
y='sepal_length_(cm)')  
ax.set_xlabel('Id')  
ax.set_ylabel('Sepal length (cm)')
```

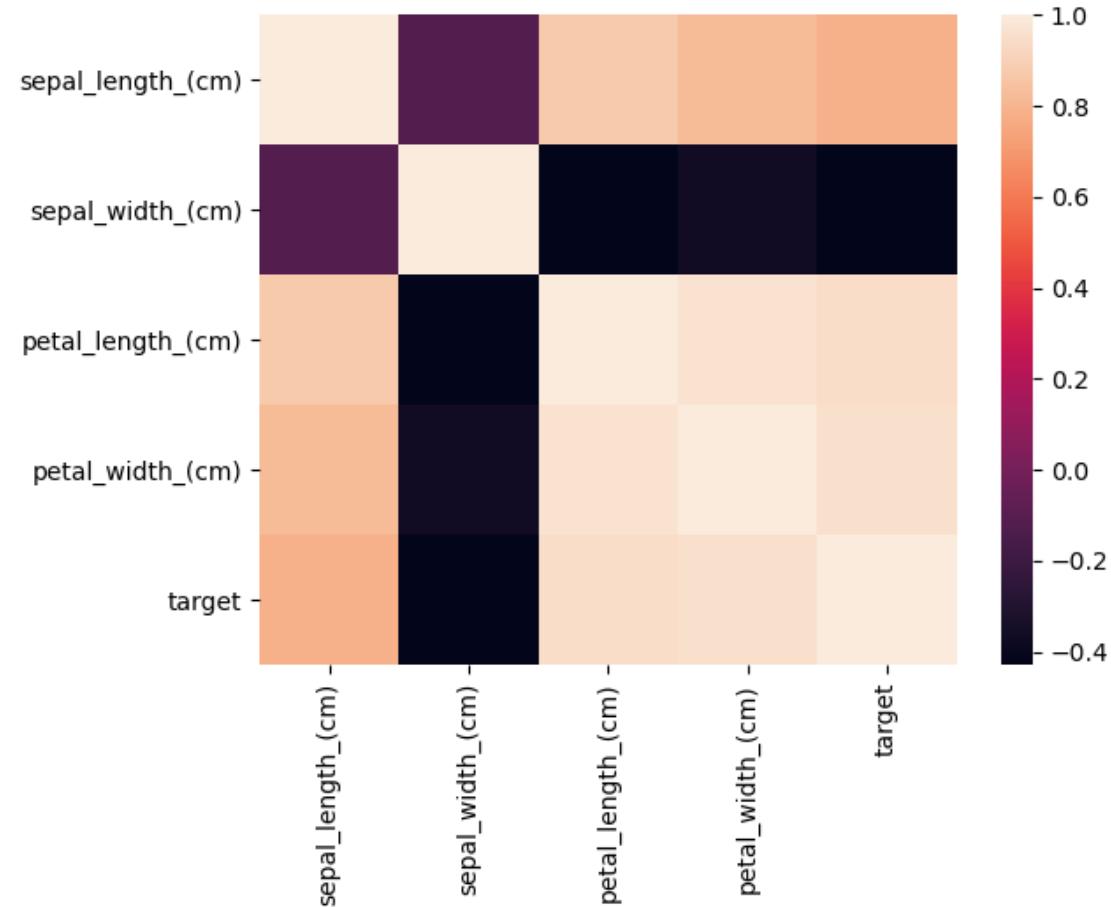




# Heatmap (B) – Seaborn

- Heat Map is commonly used to visualize the density, frequency, correlation, or intensity of data – takes as input 2 features (2D data in general)
  - Example: heat map visualizes the correlation between numerical features of the dataset

```
fig, ax = plt.subplots()  
sns.heatmap(data=df2.corr())
```





# Correlation

---

- Correlation method measures the linear relationship between 2 vars
- The correlation coefficient can never be less than -1 or higher than +1
  - +1 = there is a perfect linear relationship between the variables (as one variable increases, the other increases proportionally)
  - 0 = there is no linear relationship between the variables
  - -1 = there is a perfect negative linear relationship between the variables
    - In case of 2 highly correlated features, one of them can be removed from dataset prior running into machine learning algorithms so as to make the learning algorithm faster
      - curse of dimensionality: less features usually mean high improvement in terms of speed
    - If the dataset is small (low number of features) and speed is not an issue, perhaps don't remove these features right away. If you have correlated features but they are also correlated to the target (if target is numerical), you want to keep them
      - Some algorithms like Naive Bayes actually directly benefit from "positive" correlated features. And others like random forest may indirectly benefit from them.
    - Moral of the story, removing these features might be necessary due to speed, but remember that you might make your model worse

# Hands on

---



- Download Lab03\_Visualizations.ipynb from Piazza and follow the guidelines
-