

DSC510: Introduction to Data Science and Analytics

Lab 4: Data Preparation

Pavlos Antoniou
Office: B109, FST01



University of Cyprus
Department of
Computer Science

Prepare data for machine learning



- Data preparation is the process of making raw data ready for analysis and predictive modeling. It involves three main steps:
 - **1. Gather & Combine**
 - Collect data from databases, files (.csv, .json), APIs, social media, or statistical services.
 - Combine datasets (e.g., add features from multiple sources) to enrich the information for modeling.
 - **2. Preprocess**
 - **Clean:** fix, delete, or fill missing values.
 - **Encode:** convert categorical/text data into numerical form.
 - **Resample:** adjust the frequency of observations (important for time-series).
 - **3. Transform**
 - **Scale:** normalize or standardize features.
 - **Unskew:** adjust distributions to be more symmetrical.
 - **Feature Engineering:** select or extract features to improve model performance (covered in the next lab).



Cleaning data

- Fixing Data Formats
 - Mixed or inconsistent formats
 - Data from international sources may have varying formats that don't match expected numeric or date conventions.
 - Numeric issues
 - Decimal separator may be a comma → convert to dot (.)
 - Thousands separators may exist → remove them
 - Monetary symbols before or after numbers → strip them
 - Date/time issues
 - Dates may appear as integers (e.g., 20090609231247) instead of ISO 8601 (2009-06-09 23:12:47)
 - Transformation is needed to match Python's expected formats for plotting and analysis



Cleaning data

- **Deleting Missing Values**
 - When to delete rows:
 - If the number of missing-value rows is small relative to the dataset
 - Example: less than 10% of all rows
 - If the rows do not contain important information
 - Example: missing a non-critical category
 - When NOT to delete rows:
 - Missing values may carry important information
 - Example: missing ages could correspond to older or privacy-conscious users, relevant for decision-making
 - Key point:
 - Deleting missing values is not trivial, especially for unfamiliar datasets
 - Implementation:
 - Can be done in Python using `pandas.dropna()` (see next slides)



Cleaning data

- **Filling Missing Values**
 - Usually done column by column
 - Categorical data (e.g., device type, country): add a new category “*Unknown*”
 - Numerical data (e.g., age):
 - Use mean/median on all values of the column
 - or aggregation (e.g. mean) on similar category values
 - Time-series data: use interpolation (see sampling slides)
 - Predictive imputation: build a model to predict missing values
- **Correcting Erroneous Values**
 - Detect outliers/invalid values with [statistical analysis](#) or visualization (e.g., box plots)
 - Examples:
 - Non-numeric value in a numeric column
 - Age < 0 or > 100
 - Strategy: delete or recode → then treat as missing values

Cleaning data



- **Handling Outliers** (values that differ significantly from the rest)
 - Can distort models (e.g., regression, k-NN – especially distance-based models)
 - Removing them should be done carefully → sometimes they carry meaningful information
 - See [Appendix](#) for outlier removal techniques
- **Standardizing Categories**
 - User-entered data often contains inconsistencies:
 - Spelling mistakes, language differences, multiple formats
 - Examples:
 - Country → USA, United States, U.S.
 - Dates → 1982-10-01, 1/10/1982
 - Goal: ensure one consistent version per value



Missing values manipulation

- Methods to deal with missing values in the data frame:

df.method()	description
dropna()	Drops observations (rows) with at least one missing value
dropna(axis=1)	Drops the columns where at least one value is missing
dropna(thresh = 5)	Drops rows that contain less than 5 non-missing values
fillna(0)	Replaces missing values with a specified value
ffill()	Replace missing values by propagating the last valid observation to next valid
bfill()	Replace missing values by using the next valid observation to fill the gap
isnull()	returns True if the value is missing
notnull()	Returns True for non-missing values



Missing values manipulation

- Evaluate the number of missing values per column:

```
In [ ]: # isnull() converts missing values to True, False otherwise; sum()
          works on Boolean values, treating True=1 and False=0, counts 1s
missing_values_count = nfl_data.isnull().sum()

# look at the # of missing points in the first ten columns
missing_values_count[0:10]
```

```
Out[ ]: Date          0
        GameID        0
        Drive         0
        qtr           0
        down         54218
        time          188
        TimeUnder     0
        TimeSecs      188
        PlayTimeDiff   374
        SideoffField   450
        dtype: int64
```

- The more missing values a feature (column) has, the less reliable the data in that column might be → feature is less important
- Knowing how many missing values exist helps assess whether to keep, impute, or drop those columns
 - If a column or row has too many missing values, you might consider dropping it from the analysis, as it could contribute little to the model's accuracy



Drop columns with missing values

```
In [ ]: # remove all columns with at least one missing value  
columns_cleaned = nfl_data.dropna(axis=1)  
columns_cleaned.head()
```

```
Out[ ]:      Date      GameID  Drive  ...  ExPoint_Prob  TwoPoint_Prob  Season  
0  2009-09-10  2009091000      1  ...          0              0      2009  
1  2009-09-10  2009091000      1  ...          0              0      2009  
2  2009-09-10  2009091000      1  ...          0              0      2009  
3  2009-09-10  2009091000      1  ...          0              0      2009  
4  2009-09-10  2009091000      1  ...          0              0      2009  
  
[5 rows x 41 columns]
```

```
In [ ]: # just how much data did we lose?  
print("Columns in original dataset: %d" % nfl_data.shape[1])  
print("Columns with na's dropped: %d" % columns_cleaned.shape[1])
```

```
Out[ ]: Columns in original dataset: 102  
Columns with na's dropped: 41
```

We can specify which columns to check for missing values and drop only the rows where those columns contain NaN

```
nfl_data = nfl_data.dropna(subset=["down", "SideofField"])
```



Fill in missing values

- One option we have is to specify what we want the NaN values to be replaced with

```
In [ ]: # replace all NA's with 0
nfl_data = nfl_data.fillna(0)
# replace all NA's with 0 for a specific column
nfl_data['yacWPA'] = nfl_data['yacWPA'].fillna(0)
```

- Another option is to replace missing values with the first valid value comes after it in the same column
 - This makes a lot of sense for datasets where the observations have some sort of logical order

```
In [ ]: # replace all NA's the first valid value that comes after it in the
        # same column, then replace all the remaining na's (if any) with 0
nfl_data = nfl_data.bfill(axis=0).fillna(0)
```



Fill in missing values with imputation

- Imputation fills in the missing value with some number
- Imputed value won't be exactly right in most cases, but it usually gives more accurate models than dropping the column entirely

```
In [ ]: # Using Sklearn's simple imputer
from sklearn.impute import SimpleImputer
import numpy as np
my_imputer = SimpleImputer(missing_values=np.NaN, strategy='mean')
nfl_data[['yacWPA']] = my_imputer.fit_transform(nfl_data[['yacWPA']])
```

DataFrame nfl_data[['yacWPA']] accepted.
Series nfl_data['yacWPA'] not accepted



- SimpleImputer takes two arguments such as missing_values and strategy
 - Strategy can be set to mean, median, most_frequent, constant (with fill_value argument)
 - Numerical missing values: mean, median, most frequent, constant
 - Categorical missing values: most frequent, constant
 - fit_transform method is invoked on the instance of SimpleImputer to impute the missing values



Fill in missing values with imputation

- Strategy = mean

```
Date      GameID    ...          yacWPA  Season
0 2009-09-10 2009091000 ...        NaN     2009
1 2009-09-10 2009091000 ...  0.03689896441538476 2009
2 2009-09-10 2009091000 ...        NaN     2009
3 2009-09-10 2009091000 ... -0.1562385319864913 2009
4 2009-09-10 2009091000 ...        NaN     2009
```



Before imputation

`nfl_data.head()`

```
Date      GameID    ...          yacWPA  Season
0 2009-09-10 2009091000 ... -0.010492 2009
1 2009-09-10 2009091000 ...  0.03689896441538476 2009
2 2009-09-10 2009091000 ... -0.010492 2009
3 2009-09-10 2009091000 ... -0.1562385319864913 2009
4 2009-09-10 2009091000 ... -0.010492 2009
```

After imputation



Encoding categorical data

- Machine learning models require all features to be numerical
- Categorical text-based data (e.g. a column with “male”, “female” values) must be encoded to numbers
- Popular techniques:
 - Label Encoding
 - Ordinal Encoding
 - One-Hot Encoding (or Dummy Variable Encoding)
 - Effect Encoding
 - Bin counting
 - Feature Hashing
- Scikit-learn lib involves a few encoders but **category_encoders** lib has more with useful properties

```
conda install -c conda-forge category_encoders  
(website: http://contrib.scikit-learn.org/category\_encoders)
```



Label Encoding

- Assigns a unique integer value to each category of a categorical feature
 - e.g. "Red" → 0, "Green" → 1, "Blue" → 2
- Used for **nominal** (unordered) **categorical features**
 - numbers assigned are arbitrary - don't represent ranking or size
- Potential issue: implies ordinal relationships between categories
 - e.g. Red (0) seems to be closer to Green (1) than to Blue (2)
 - high ordinal values possess higher “weight” and may be considered of higher importance especially in distance-based ML techniques



Ordinal Encoding

- Similar to label encoding, but respects category order
- Used for **ordinal categorical features** (with natural, ordered values) and when **retaining the order is important**
- Encoding reflects the sequence

Degree
0 High school
1 Masters
2 Diploma
3 Bachelors
4 Bachelors
5 Masters
6 Phd
7 High school
8 High school

Natural order:
'High school':1
'Diploma':2
'Bachelors':3
'Masters':4
'Phd':5

Label & Ordinal Encoding: OrdinalEncoder



```
import category_encoders as ce
import pandas as pd
df=pd.DataFrame(
{'Degree':['High school', 'Masters', 'Diploma',
'Bachelors', 'Bachelors', 'Masters', 'Phd', 'High
school', 'High school']})

#Original data
df
```

	Degree
0	High school
1	Masters
2	Diploma
3	Bachelors
4	Bachelors
5	Masters
6	Phd
7	High school
8	High school

```
# create object of Ordinal encoding
ordinal_encoder = ce.OrdinalEncoder(
mapping=[{'col':'Degree','mapping':{'None':0,'High
school':1,'Diploma':2,'Bachelors':3,'Masters':4,
'Phd':5}}])

#fit and transform data
df['Ordinal'] =
ordinal_encoder.fit_transform(df['Degree'])
df
```

	Degree	Ordinal
0	High school	1
1	Masters	4
2	Diploma	2
3	Bachelors	3
4	Bachelors	3
5	Masters	4
6	Phd	5
7	High school	1
8	High school	1

Note: If no mapping is given, order is automatically chosen by the encoder → **Label encoding**



One-Hot Encoding

- Used for **nominal** (unordered) **categorical features**.
 - Prevents imposing false orderings (as happens with label encoding).
 - Creates a new binary feature (column) for each category value.
 - 1 = category present
 - 0 = category absent
 - These binary features are called dummy variables.
 - Number of dummy variables = number of categories in the original feature.
 - Preferred over label encoding for distance-based models.
-



One-Hot Encoding: OneHotEncoder

```
# Create object for One-hot encoding  
onehot_encoder=ce.OneHotEncoder(cols=['Degree'], use_cat_names=True)  
#fit and transform data  
df_onehot = onehot_encoder.fit_transform(df)  
df_onehot
```

a list of columns to encode, if None, all string columns will be encoded

Dummy variables

	Degree	Ordinal
0	High school	1
1	Masters	4
2	Diploma	2
3	Bachelors	3
4	Bachelors	3
5	Masters	4
6	Phd	5
7	High school	1
8	High school	1



	Degree_High school	Degree_Masters	Degree_Diploma	Degree_Bachelors	Degree_PhD	Ordinal
0	1.0	0.0	0.0	0.0	0.0	1
1	0.0	1.0	0.0	0.0	0.0	4
2	0.0	0.0	1.0	0.0	0.0	2
3	0.0	0.0	0.0	1.0	0.0	3
4	0.0	0.0	0.0	0.0	1.0	3
5	0.0	1.0	0.0	0.0	0.0	4
6	0.0	0.0	0.0	0.0	1.0	5
7	1.0	0.0	0.0	0.0	0.0	1
8	1.0	0.0	0.0	0.0	0.0	1



Drawbacks of One-Hot Encoding

- Creates **one new column per category** → many columns for high-cardinality features (with high number of unique values)
 - Results in **high-dimensional datasets** (with very high number of features) when multiple categorical features are present
 - Can **slow model training** and **hurt performance** due to increased computational cost
-



Cyclical feature encoding

- For time-dependent features (e.g., month, day, hour), standard numeric encoding can be misleading.
 - Example: in the hour feature, 0 (midnight) and 23 (11 PM) are numerically far apart, but in reality they are adjacent in time.
- Step 1: Extract time components from datetime (month, day, hour, minute, weekday, etc.).
- Step 2: Apply cyclical encoding → map each value onto the unit circle using sine and cosine: $x = \sin(2\pi * \text{value} / \text{max_value})$, $y = \cos(2\pi * \text{value} / \text{max_value})$
- This encoding preserves the cyclical nature of time (e.g., hours, days, months).

	datetime	temperature	hour
9	2012-10-01 21:00:00	12.776627	21
10	2012-10-01 22:00:00	12.789767	22
11	2012-10-01 23:00:00	12.802906	23
12	2012-10-02 00:00:00	12.816046	0
13	2012-10-02 01:00:00	12.829185	1

```
data['hour_sin'] = np.sin(2 * np.pi * data['hour']/23.0)
data['hour_cos'] = np.cos(2 * np.pi * data['hour']/23.0)
```

	datetime	temperature	hour	hour_sin	hour_cos
10	2012-10-01 22:00:00	12.789767	22	-2.697968e-01	0.962917
11	2012-10-01 23:00:00	12.802906	23	-2.449294e-16	1.000000
12	2012-10-02 00:00:00	12.816046	0	0.000000e+00	1.000000
13	2012-10-02 01:00:00	12.829185	1	2.697968e-01	0.962917

11 PM is close to 12 midnight in terms of sin and cos



Data Transformation: Scaling data

- Feature rescaling
 - Some classification/regression/clustering techniques (see next slide) use the notion of distance (e.g. Euclidean) to measure similarity between 2 observations
 - Example
 - Classify houses with 2 features
 - $x_1 = \text{size} (0 - 2000m^2)$
 - $x_2 = \text{number of bedrooms} (1 - 5)$
 - $\text{Euclidean distance}(X_1, X_2) = \sqrt{(523 - 127)^2 + (4 - 2)^2}$
 - Distance is governed by features having boarder range of values
 - When distance is used by algorithms **make sure features are on a similar scale**
 - Target variable is not necessary to be scaled

$$X = \begin{bmatrix} 523 & 4 \\ 127 & 2 \\ 25 & 1 \end{bmatrix}$$

Feature x_1 with high magnitudes weights a lot more (dominates) in the distance calculations than feature x_2 with lower magnitudes



Data Transformation: Scaling data

- Some examples of algorithms where feature scaling matters are:
 - k-nearest neighbors (**kNN**) for classification uses Euclidean distance
 - k-means for clustering uses Euclidean distance
 - gradient descent/ascent-based optimization used in **logistic regression**, **Support Vector Machines (SVMs)**, neural networks etc.
 - Weights for features with higher magnitudes will update much faster than others
 - linear discriminant analysis (**LDA**), principal component analysis (**PCA**)
 - you want to find directions of maximizing the variance (under the constraints that those directions/eigenvectors/principal components are orthogonal)
- Decision trees and ensembles of trees are unaffected by the scale of feature variables. Examples:
 - bagging like RandomForest
 - boosting like AdaBoost, Gradient Boosting, XGBoost, LightGBM, CatBoost



Data Transformation: Scaling data

- Feature normalization
 - Rescales **each feature** individually into a given range, e.g. [0, 1]

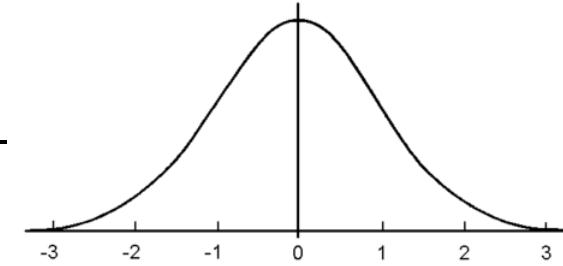
$$x = \begin{bmatrix} 4 & 13 \\ 3 & 2 \\ 7 & 8 \end{bmatrix}, x_{i,j,desc} = \frac{x_{i,j} - \min(x_j)}{\max(x_j) - \min(x_j)} \Rightarrow x_{desc} = \begin{bmatrix} 0.25 & 1 \\ 0 & 0 \\ 1 & 0.55 \end{bmatrix}$$

- Scikit-learn module: [MinMaxScaler](#) or [MaxAbsScaler](#)
 - MinMaxScaler: Transforms features by scaling each feature to a given range ($x_{\min} \rightarrow x_{\max}$).

```
from sklearn.preprocessing import MinMaxScaler
df = pd.DataFrame({'A': [4, 3, 7], 'B': [13, 2, 8] })
# create the scaler object
scaler = MinMaxScaler(feature_range=(0, 1))
# train the scaler (find min and max)
scaler.fit(df)
# scale the dataset (apply the transformation)
minMaxRescaledX = scaler.transform(df)
print(minMaxRescaledX)
```

minMaxRescaledX =
scaler.fit_transform(df)

Data Transformation: Scaling



- Feature standardization
 - Rescales **each feature individually** to make values have zero mean ($\mu = 0$) and unit variance ($\sigma^2 = 1$)
$$x = \begin{bmatrix} 4 & 13 \\ 3 & 2 \\ 7 & 8 \end{bmatrix}, x_{i,j, std} = \frac{x_{i,j} - \text{mean}(x_j)}{\sigma} \Rightarrow x_{std} = \begin{bmatrix} -0.39 & 1.86 \\ -0.98 & -1.226 \\ 1.37 & 0.07 \end{bmatrix}$$
 - Centers values around zero and adjusts their spread so that variance is 1
 - Benefits features that are approximately normally distributed (gaussian)
 - Useful for distance-based algorithms such as the SVM (RBF kernel) and when using gradient-based optimization methods
 - Scikit-learn module: [StandardScaler](#)

```
from sklearn.preprocessing import StandardScaler
df = pd.DataFrame({'A': [4, 3, 7], 'B': [13, 2, 8] })
scaler = StandardScaler()
# train the standardizer (find mean, std) and standardize the dataset
standardRescaledX = scaler.fit_transform(df)
print(standardRescaledX)
```



Data Transformation: Scaling data

- Feature robust standardization
 - When data contains outliers, mean value and variance used by the Standard Scaler can distort the rescaled values
 - MinMaxScaler is also sensitive to the presence of outliers as well
 - Robust standardization is to rescale **each feature individually** to make values have zero median (median=0) and unit interquartile range (IQR=1)
 - Centers values around 25th and 75th percentiles (within the IQR)
 - Benefits features with non-gaussian distributions, particularly those with outliers or skewed (long-tailed) distributions
 - Scikit-learn module: [RobustScaler](#)

```
from sklearn.preprocessing import RobustScaler
df = pd.DataFrame({'A': [4, 3, 7], 'B': [13, 2, 8] })
rscaler = RobustScaler().fit(df)
# train the standardizer (find median, quantiles) and standardize the dataset
robustRescaledX = rscaler.fit_transform(df)
print(robustRescaledX)
```

[Compare the effect of different scalers on data with outliers](#)

When to normalize or standardize features?



- When features have different scales (e.g., age in years vs. income in thousands)
 &
 distance-based algorithms (e.g., k-NN, SVM) or gradient-based algorithms (e.g., logistic/linear regression) will be used
 - Normalization (MinMaxScaler):
 - Works best for features with different ranges and few or no significant outliers
 - Standardization (StandardScaler):
 - Works well for features with approximately normal distributions and few/no influential outliers
 - Robust Scaling (RobustScaler):
 - Best for features with many outliers or skewed (long-tailed) distributions
-

Scale or normalize label / ordinal encoded data when using distance-based algorithms?



- Depends on the nature of your data
- When to scale:
 - Encoded feature has many levels → prevents it from dominating
 - Ordinal encoding with meaningful, evenly spaced values (e.g., 1-5 survey ratings) → can use scaling that preserves relationships (e.g. MinMax scaling)
- When not to scale:
 - Purely categorical / arbitrary values (e.g., education levels 1-3, colors 0-2)
 - Scaling such values (getting them closer or further) can imply a relationship that doesn't exist

Scale or normalize one-hot encoded data when using distance-based algorithms?



- Scaling features being one-hot-encoded **is not recommended**
 - Binary values don't require scaling
 - One-hot encoding produces binary (0/1) columns. These already express presence (1) or absence (0) clearly. Scaling adds no meaningful information.
 - Scaling distorts categorical meaning
 - Scaling would turn 0/1 values into continuous values (e.g., 0.5) that have no valid categorical interpretation.
-



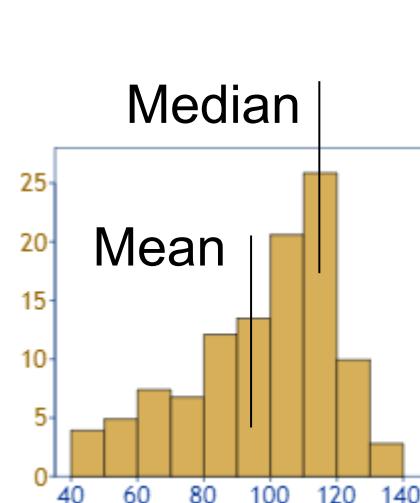
When to normalize or standardize target?

- Usually **not required**, but can help in regression tasks
- When to scale:
 - Target has a wide value range (e.g., income in hundreds to millions) → prevents large values from dominating
 - Using gradient descent-based models (e.g., linear regression) → speeds up convergence

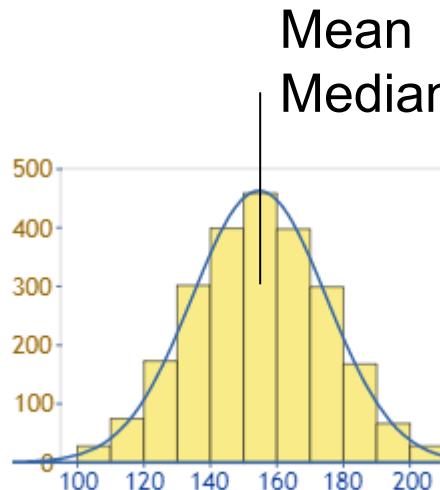


Data Transformation: Unskewing data

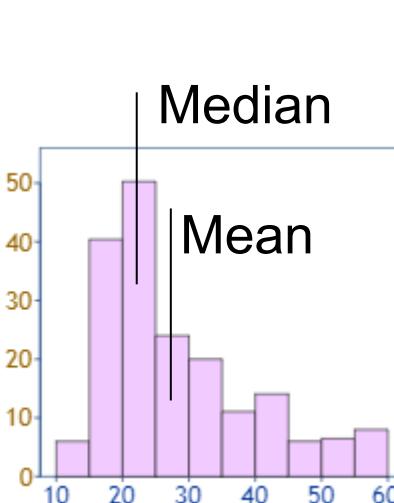
- A variable is skewed when its distribution curve is asymmetrical as compared to a normal distribution curve that is perfectly symmetrical
- *Skewness* is the measure of the asymmetry
 - The skewness for a gaussian or normal distribution is 0



Left / negative skew:
Long tail is on the left / negative side of the peak



No skew (symmetrical distribution)

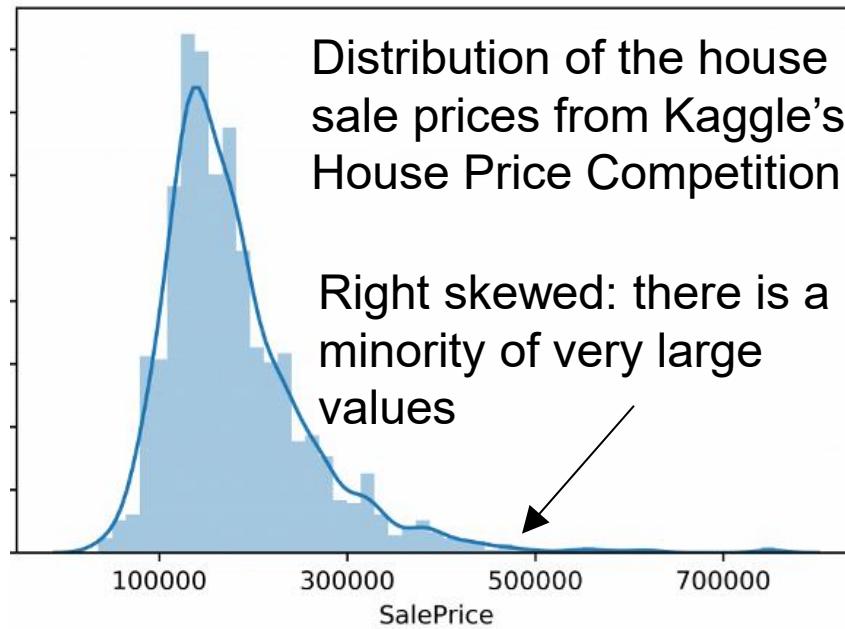


Right / positive skew:
Long tail is on the right / positive side of the peak



Effects of skewed data

- Skewness of (an input or target) variable may degrade the predictive model's ability to predict values towards the long tail side



- A regression model for predicting house sale prices or using house sale price as input feature (using the above dataset) will be trained on a much larger number of moderately priced houses and will be less likely to successfully predict the price for the most expensive houses



Unskewing transformations

- **Goal:** Transform skewed data to be more symmetric (Gaussian)
- **Why:** Some models (Linear/Logistic Regression, SVM, Gaussian NB) work better with near-Gaussian data
- **Not needed for:** Tree-based models (Decision Trees, Random Forests, Boosting) → not affected by skewness
- **Where:** Apply on highly skewed features or target
- Scalers (MinMax, Standard, Robust) do not change the skew (shape) of the distribution – use other transformations (next slide)



Unskewing transformations

- **Square Root (SQRT) transformation**

```
import numpy as np  
np.sqrt(df.column)
```

- **Log(arithmic) transformation**

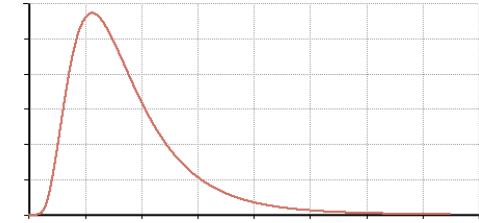
```
import numpy as np  
np.log(df.column)
```

- **Boxcox & Yeo-Johnson transformations**

- **Box-Cox** can handle both right and left skewed distributions but can only be applied to values that are **strictly positive**

- **Yeo-Johnson** can also handle both right and left skewed distributions and can be applied to both positive and negative values

- work well on right skewed distributions



- applicable on features with strictly positive values (sqrt and log cannot be applied on negative values)

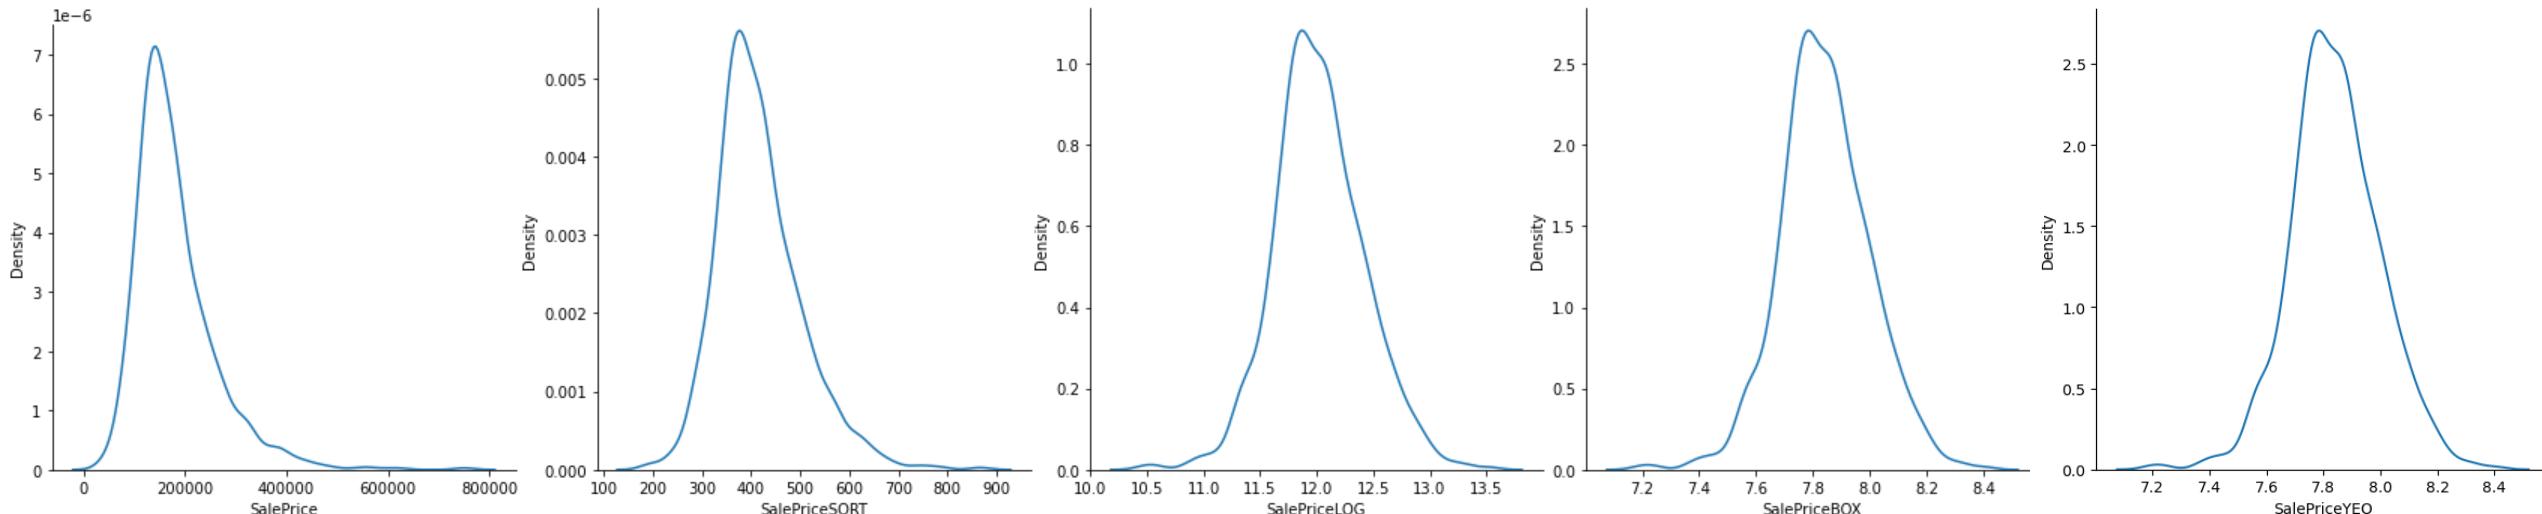
```
from scipy.stats import boxcox  
df['bc_col'] = boxcox(df['col'])
```

```
from scipy.stats import yeojohnson  
df['yj_col'] = yeojohnson (df['col'])
```



Unskewing transformations: Examples

- Pandas .skew() method can be used to measure skewness of data



Original SalePrice column
Skewness: 1.8828757

SQRT transformation
Skewness: 0.9431527

LOG transformation
Skewness: 0.1213351

BoxCox transformation
Skewness: -0.0086530

YeoJohnson transform
Skewness: -0.0086536

- Source code and results are available in .ipynb file in course website
- A quite descriptive document on skewness can be found [here](#)



APPENDIX



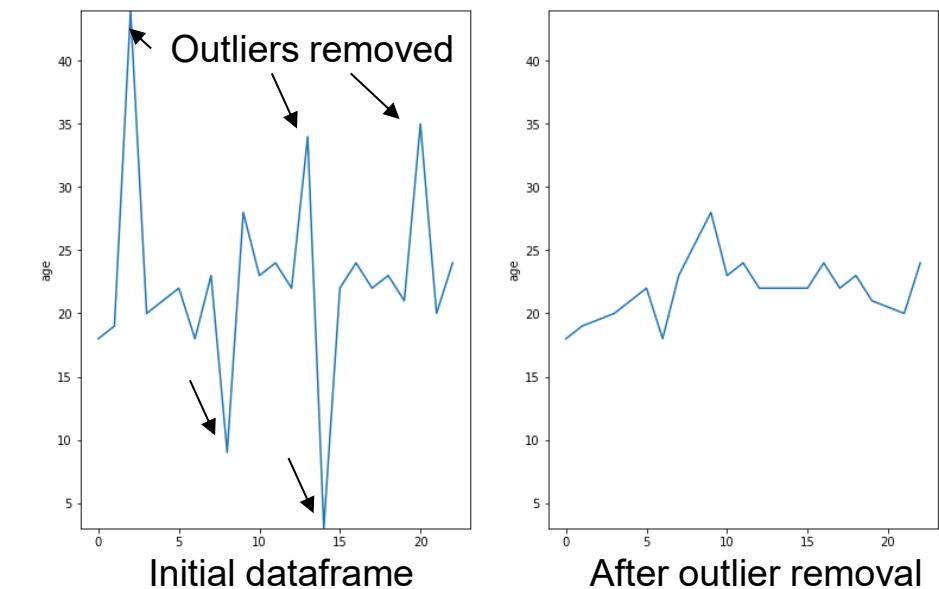
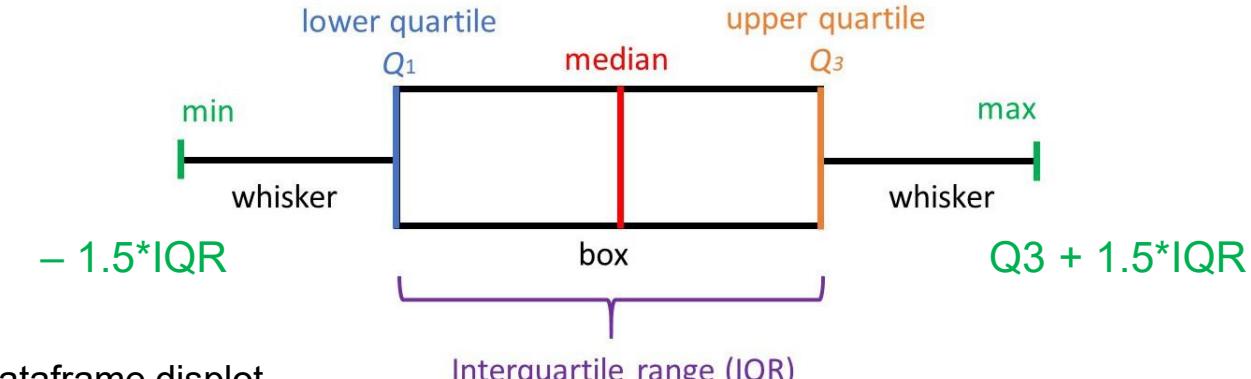
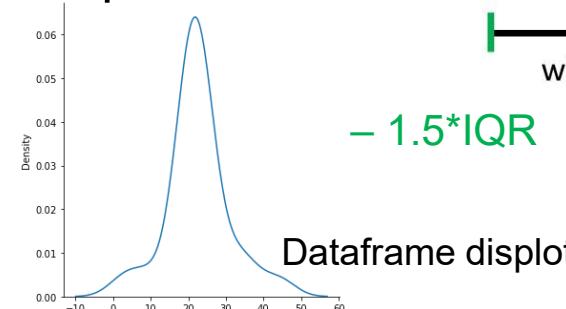
Removing outliers

- Methods for removing outliers **on each feature independently:**

- Interquartile Range (IQR) method

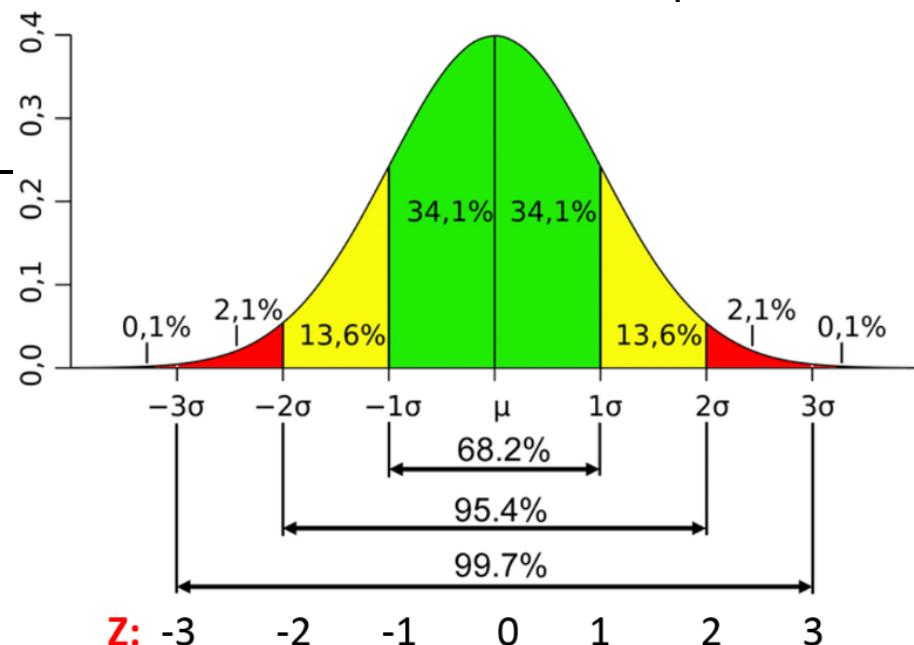
- Outliers are considered data points:
 - below $Q_1 - 1.5 \cdot IQR$
 - above $Q_3 + 1.5 \cdot IQR$

```
import seaborn as sns
import matplotlib.pyplot as plt
df2 = pd.DataFrame({'age': [18, 19, 44, 20, 21, 22, 18, 23, 9, 28, 23, 24, 22, 34, 3, 22, 24, 22, 23, 21, 35, 20, 24]})
plt.subplot(1,2,1)
sns.lineplot(data=df2, y=df2['age'], x=df2.index)
plt.ylim([df2['age'].min(), df2['age'].max()])
Q1 = df2['age'].quantile(0.25)
Q3 = df2['age'].quantile(0.75)
IQR = Q3-Q1
maximum = Q3 + 1.5*IQR
minimum = Q1 - 1.5*IQR
df3 = df2[ (df2['age'] > minimum) & (df2['age'] < maximum) ]
plt.subplot(1,2,2)
sns.lineplot(data=df3, y=df3['age'], x=df3.index)
plt.ylim([df2['age'].min(), df2['age'].max()])
```

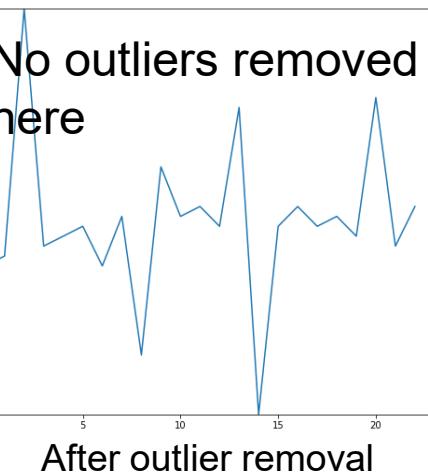
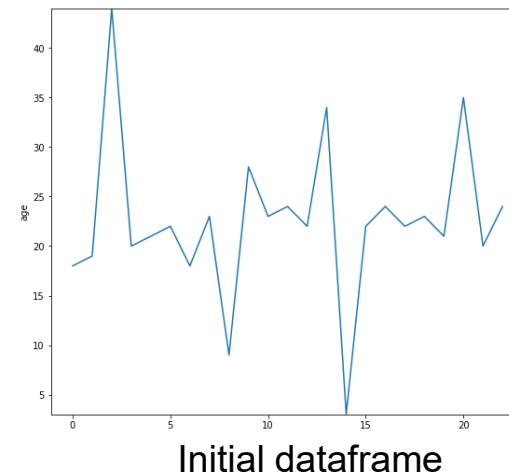


Removing outliers

- Mean (μ) and Standard Deviation (σ) method
 - For features that follow the normal distribution
 - Outliers are considered data points:
 - below $\mu - 3\sigma$
 - above $\mu + 3\sigma$



```
mean = df2['age'].mean()  
std = df2['age'].std()  
maximum = mean + 3*std  
minimum = mean - 3*std  
df4 = df2[ (df2['age'] > minimum) & (df2['age'] < maximum) ]  
plt.subplot(1,2,2)  
sns.lineplot(data=df4, y=df4['age'], x=df4.index)  
plt.ylim([df2['age'].min(), df2['age'].max()])
```





Removing outliers

– Median and Median Absolute Deviation (mad) method

- Replaces the mean and standard deviation with more robust statistics such as the median and median absolute deviation
- Outliers are considered data points:
 - below median – $3 \cdot \text{mad}$
 - above median + $3 \cdot \text{mad}$

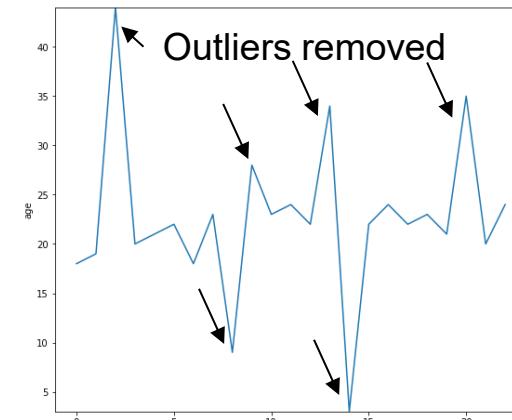
$$MAD = \text{Median}(|X_i - \text{median}|)$$

Step 1: Find the median.

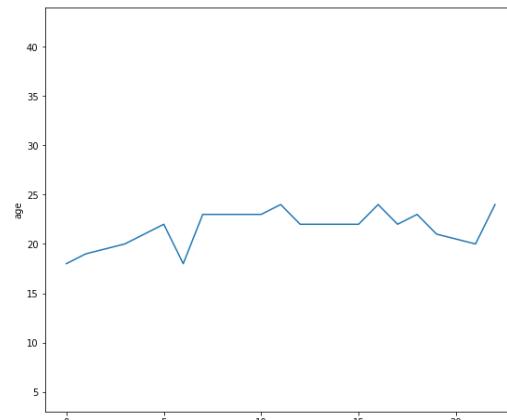
Step 2: Subtract the median from each x-value using the formula $|x_i - \text{median}|$.

Step 3: find the median of the absolute differences.

```
import scipy as sp
median = df2['age'].median()
mad = sp.stats.median_abs_deviation(df2['age'])
maximum = median + 3*mad
minimum = median - 3*mad
df4 = df2[ (df2['age'] > minimum) & (df2['age'] < maximum) ]
plt.subplot(1,2,2)
sns.lineplot(data=df4, y=df4['age'], x=df4.index)
plt.ylim([df2['age'].min(), df2['age'].max()])
```



Initial dataframe



After outlier removal



Mean/std vs Median/mad

- Mean and std are highly affected by outliers
 - All values (including outliers) are used to calculate the mean and std
- Median and MAD are not highly affected by outliers
 - Outlier changes only center value(s) which are used to calculate the median
- Example:
 - dataset: {2,3,5,6,9}
 - mean = 5, std = 2.738, median = 5, mad = 2
 - Add outlier value 1000 to dataset
 - dataset: {2, 3, 5, 6, 9, 1000}
 - mean = 170.83, std = 406.21, median = $(5+6)/2 = 5.5$, mad = 3
 - The outlier
 - increases mean by 165.83 and std by 403.472
 - increases median by 0.5 and mad by 1