

DSC510: Introduction to Data Science and Analytics

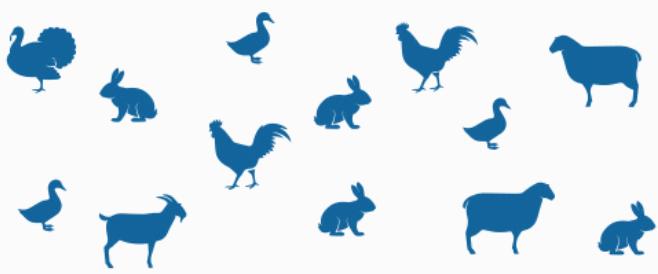
Lab 9: Clustering



University of Cyprus
Department of
Computer Science

Pavlos Antoniou

Office: B109, FST01



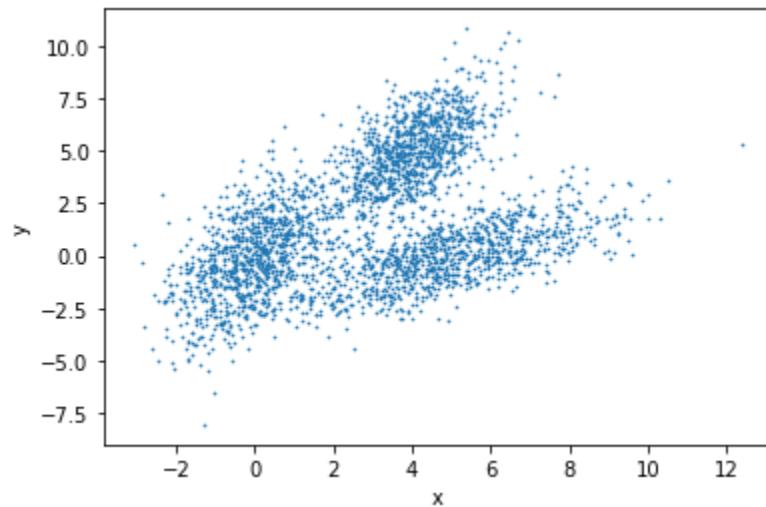
Classification

Clustering

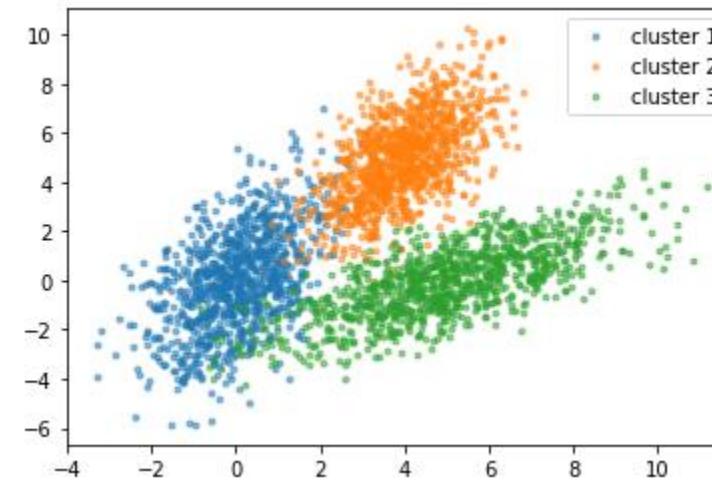
Clustering



- Unsupervised Machine Learning process of dividing the dataset into groups consisting of similar data points
- Each group is called a cluster and contains data points with high similarity and with low similarity with data points in other clusters



Samples in two-dimensional (2 features) space
BEFORE clustering



Samples in two-dimensional (2 features) space
AFTER clustering in three groups

The number of clusters (k) is a hyper parameter of clustering models: needs to be defined prior performing clustering



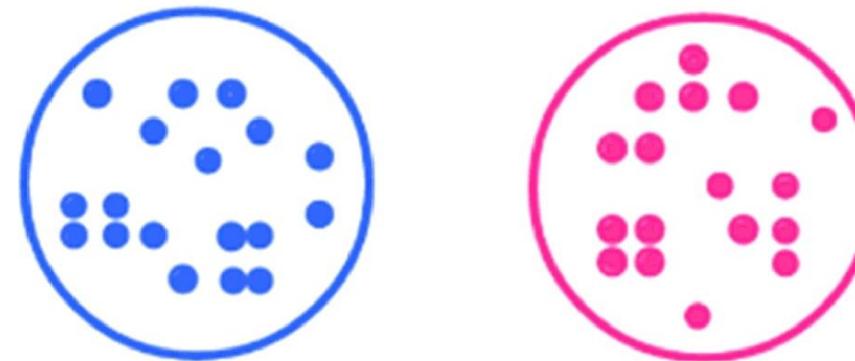
Types of Clustering

- Exclusive (non-overlapping) clustering
 - Overlapping clustering
 - Hierarchical clustering
-



Exclusive (non-overlapping) clustering

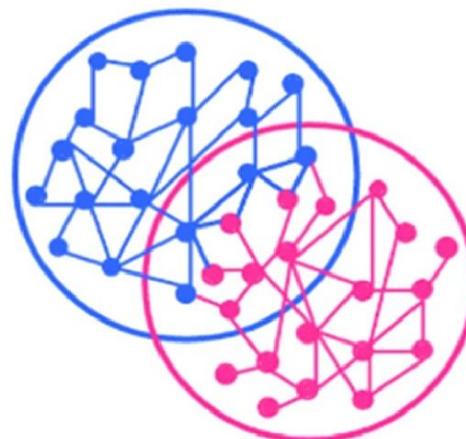
- Hard clustering
- Data point belongs exclusively to one of the identified disjoint clusters
- Example algorithm: [K-means Clustering](#)





Overlapping clustering

- Soft clustering
- Data point belongs to more than one clusters
 - data point can belong to a cluster with some degree of membership (probability) between 0 and 1
- Example algorithm: Fuzzy / C-means Clustering





Hierarchical clustering

- Hierarchy of clusters is identified using either agglomerative (bottom-up) or divisive (top-down) approaches
- Agglomerative approach: we consider each data point as one cluster and we iteratively merge them according to a criterion e.g. distance

- Create a matrix that contains all pairwise ([linkage](#)) distances among all clusters



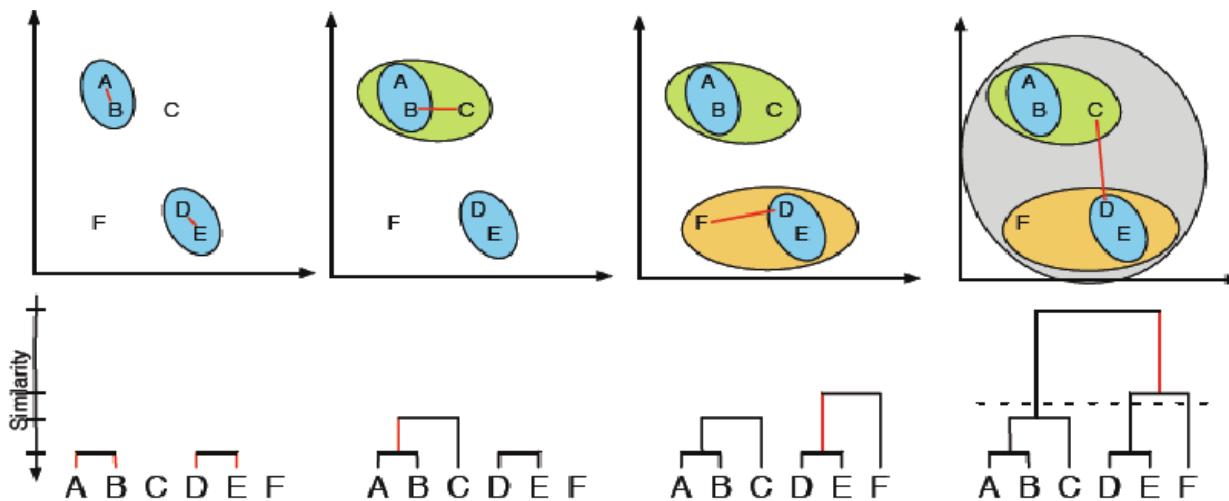
- The closest two clusters are merged

- Main output of Agglomerative clustering is a dendrogram which shows the hierarchical relationship between clusters

- The height of the vertical dendrogram lines reflects the distance (or dissimilarity) between the 2 merged clusters

- Iterative process continues until all the clusters are merged together or the desired number of clusters is reached

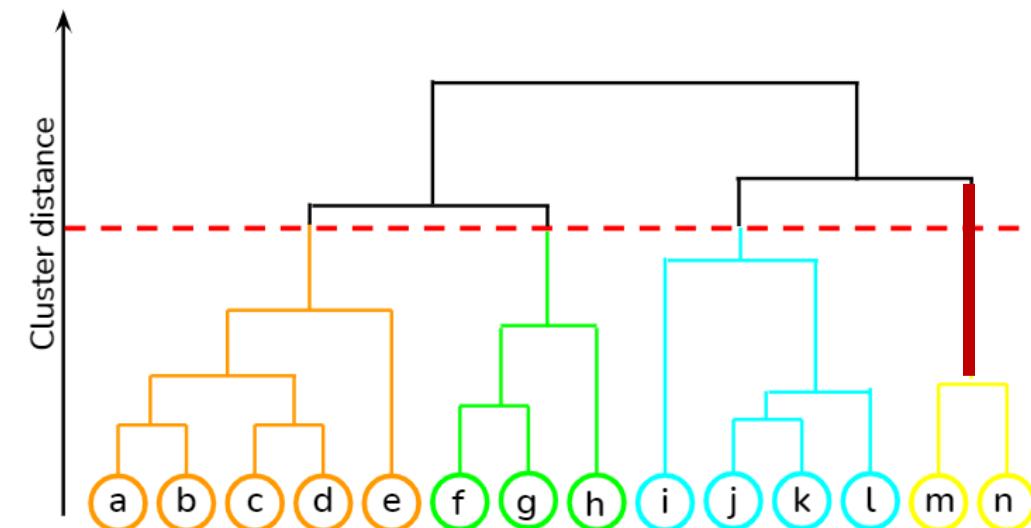
- Example algorithm: [Agglomerative clustering](#)





Which is the optimal num of clusters?

- To find the “optimal” number of clusters, you:
 - Look for the highest vertical line between successive merges (horizontal lines)
 - Draw a horizontal cut across the dendrogram at that gap.
 - The number of clusters = the number of vertical lines the cut passes through.
- Intuitively:
 - Cutting the dendrogram at a large jump in linkage distance separates groups that are significantly different
 - Cutting too low → too many small clusters
 - Cutting too high → merges dissimilar clusters





Hierarchical clustering

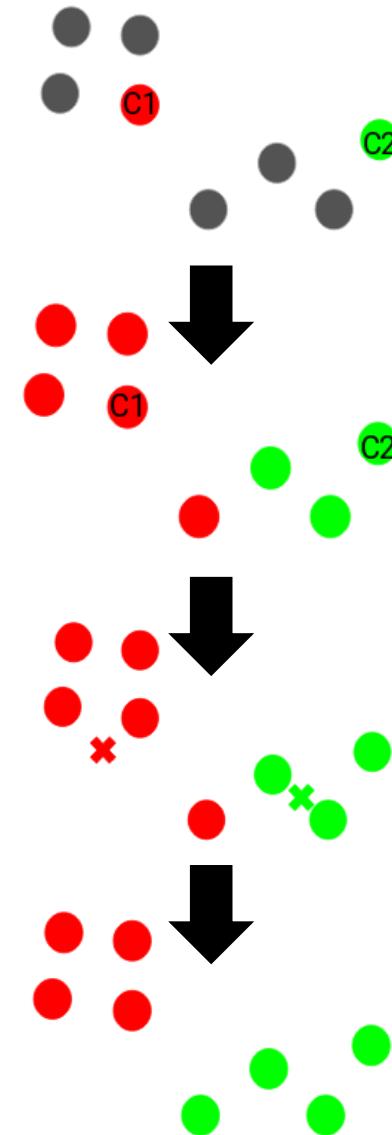
- Divisive approach: we consider the whole dataset as one cluster and iteratively split them into multiple clusters according to some evaluation criterion (e.g, distance-based criteria, statistical or variance-based criteria)
 - Example algorithm: [Bisecting K-Means](#)





K-means: exclusive clustering algorithm

- User provides the number of clusters K
- K-means iterative process involves the following steps:
 1. Selects K samples from data, or generates K points to be used as centroids
 2. Assigns all samples to the closest cluster centroid
 3. Recomputes the centroids of newly formed clusters
 - Centroid can be calculated as the mean of data points of the cluster
 - E.g. data points: (80, 56), (75, 53), (60, 50), (68, 54) Centroid: (70.75, 53.25)
 4. Repeats steps 2 and 3
- Stopping criteria for K-means:
 - Centroids of newly formed clusters do not change
 - Samples remain in the same cluster
 - Maximum number of iterations is reached





K-means Issues

- Works only with numerical data
 - Encoding techniques needed to convert categorical data to numerical
- Distance-based algorithm: Euclidean distance
 - Features should be of similar scale – data scaling needed e.g. Standard scaler, Robust scaler, etc.
- Depends on initial centroid selection
 - The more optimal the positioning of these initial centroids, the fewer iterations of the k -means algorithm will be required for convergence
 - Strategic consideration to the initialization of these initial centroids could prove useful
 - Available initialization strategies:
 - Random selection: prone to bad selection (e.g. very close to each other)
 - K-means++ is a smart centroid initialization technique which selects centroids being as far as possible from one another



K-means bottom line

- Easy to implement and use
 - Needs to scale features if in different scales
 - Good initial centroid selection method available: K-means++
 - Need to set K prior running the algorithm
-

Choosing the best K (number of clusters)

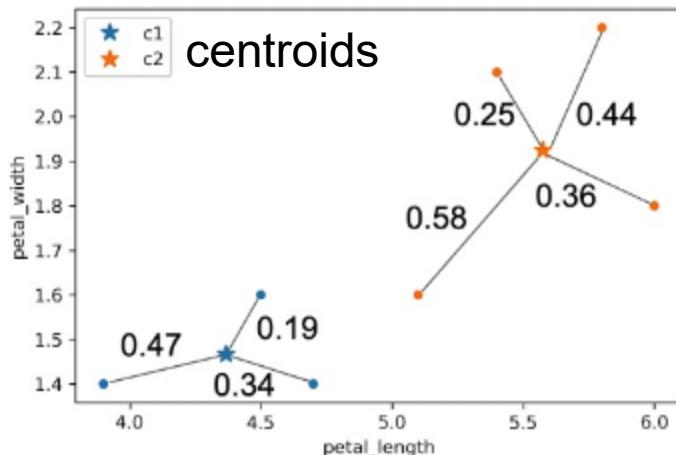


- How can we determine the “best” value of K ?
 - Is there an objective method?
- An estimation can be obtained using the following techniques:
 - Elbow method
 - Silhouette analysis



Elbow method parameters

- **Inertia:** The sum of squared distances* from each sample (data point or row) to its assigned cluster centroid
 - (*) Typically, the Euclidean distance metric is used
- **Distortion:** Weighted (by the cluster size) sum of squared distances from each sample (data point) to its assigned cluster centroid
- Example: use K-means, with K=2, to cluster 7 data points of a dataset **having only 2 features** in order to be able to visualize the distances in the 2-dimensional space and better understand the calculations below:



$$\text{Inertia} = 0.47^2 + 0.19^2 + 0.34^2 + 0.25^2 + 0.44^2 + 0.36^2 + 0.58^2$$

$$\text{Distortion} = \frac{(0.47^2 + 0.19^2 + 0.34^2)}{3} + \frac{(0.25^2 + 0.44^2 + 0.36^2 + 0.58^2)}{4}$$

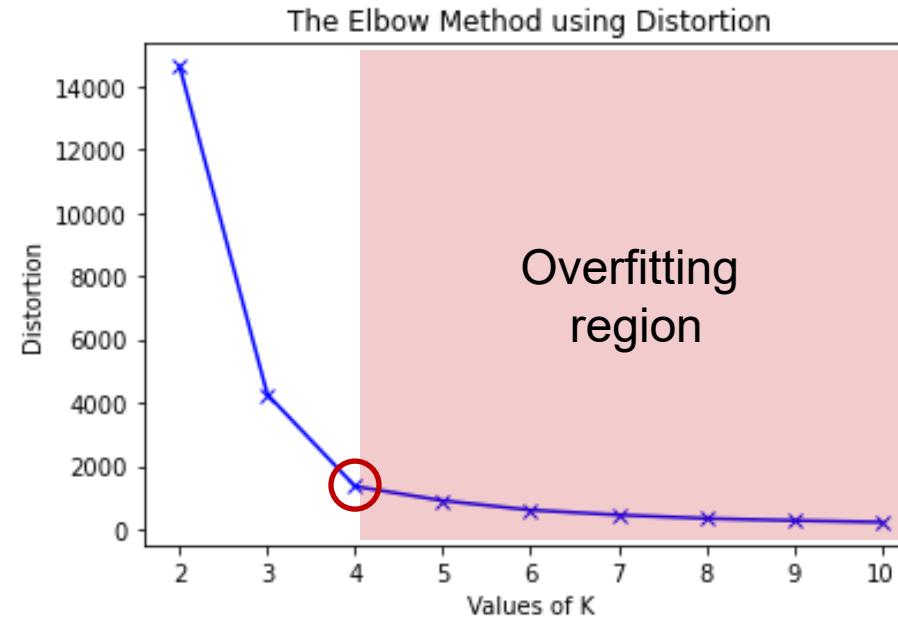
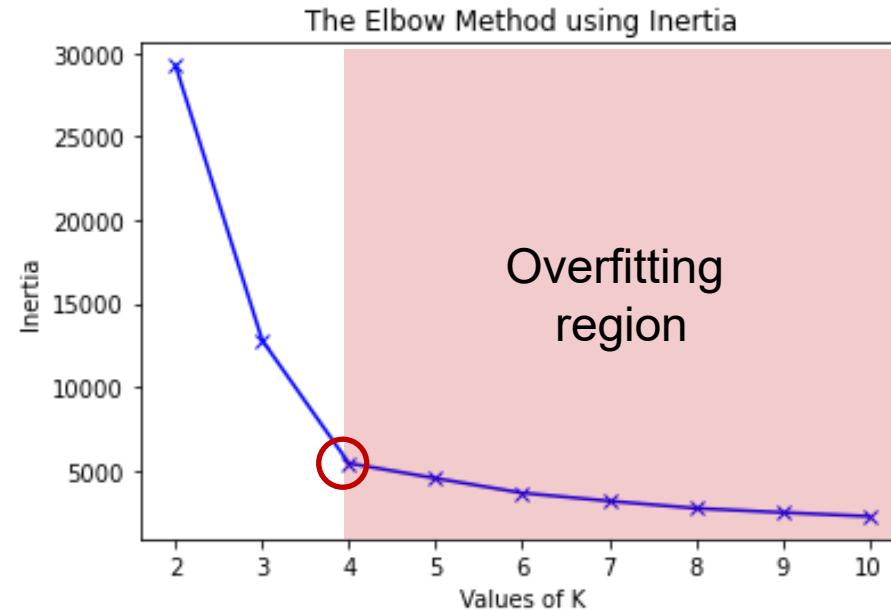


- Inertia measures:
 - Compactness (Cohesion) → how tightly grouped the data points are within each cluster.
 - Explained variance → lower inertia means the clustering explains the data better (less error).
- Inertia does *not* measure:
 - Separation between clusters (it ignores how far clusters are from each other).
 - Cluster shape or density (it assumes spherical clusters, as in k-means).
- Distortion measures:
 - The average compactness per data point.
 - How much “error” or “information loss” occurs when representing data by its cluster centers.



Elbow method using inertia / distortion

- Run K-means for different K and plot the values of inertia / distortion

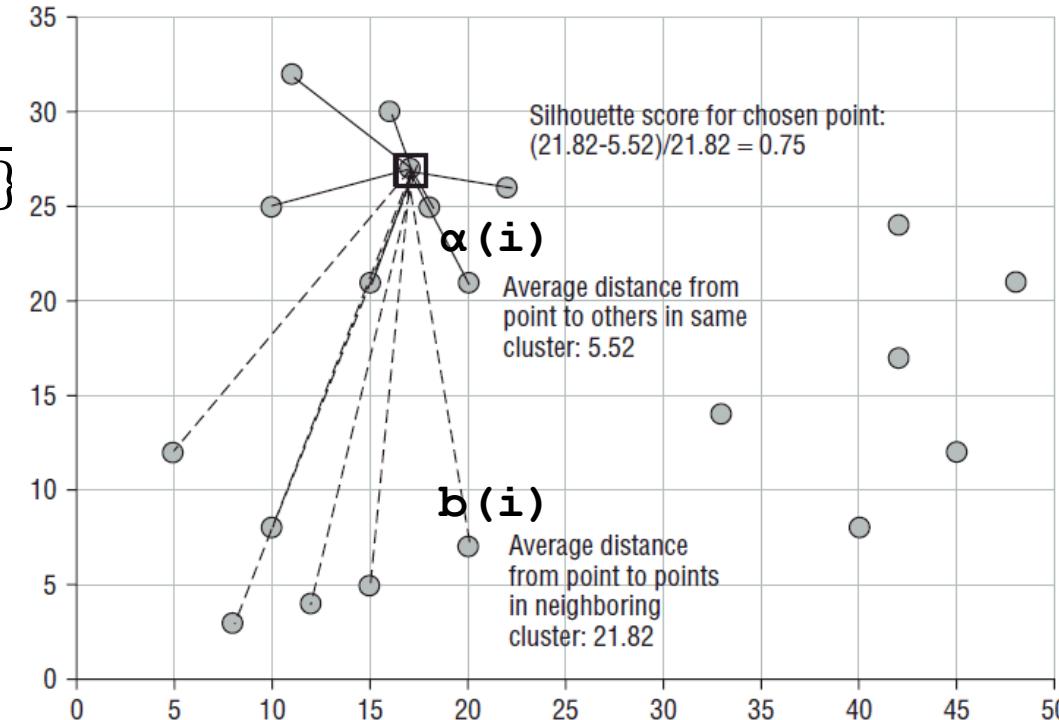


- As the number of clusters K increases inertia/distortion **always decreases** → clusters get smaller, points are closer to their centers
- After a certain K, improvement becomes marginal → “elbow” point
- That elbow corresponds to a good trade-off between **compactness (low inertia)** and **simplicity (few clusters)**



Silhouette score

- Measures how similar a data point is to its own cluster (compactness) compared to other clusters (separation)
 - Silhouette score for data point i :
$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$$
 - ranges from -1 to 1
 - high value indicates that the data point is well matched to its own cluster and poorly matched to neighboring clusters
 - values near 0 indicate overlapping clusters
 - negative values generally indicate that a sample has been assigned to the wrong cluster, as a different cluster is more similar
 - Find mean value of silhouette score of all data points => if most objects have a high value, then mean value is close to 1 and the clustering configuration is appropriate





Example: Drivers Dataset

- Includes 4000 drivers
- Each observation has 3 columns:
 - Driver_ID
 - Distance_Feature: mean distance covered per day
 - Speeding_Feature: mean percentage of time a driver was >5 mph over the speed limit
- No target variable (no notion of groups, labels)
- Load dataset, drop Driver_ID column, scale features

```
dataset = pd.read_csv('fleet_data.csv')
dataset = dataset.drop(columns=['Driver_ID'])
scaler = RobustScaler()
X = scaler.fit_transform(dataset)
```

- Source code for all clustering experimentations can be found [here](#)

K-Means



- Python implementation: `sklearn.cluster.Kmeans()` class
- Run algorithm to define groups (clusters)

```
from sklearn.cluster import Kmeans  
  
# create K-means object and run clustering on the input values (X)  
# default k (n_clusters param) → 8  
# default centroid initialization method (init param) → k-means++  
kmeans = KMeans(n_clusters=2).fit(X)  
print(kmeans.labels_) # labels of each sample  
print(kmeans.centroids_)
```

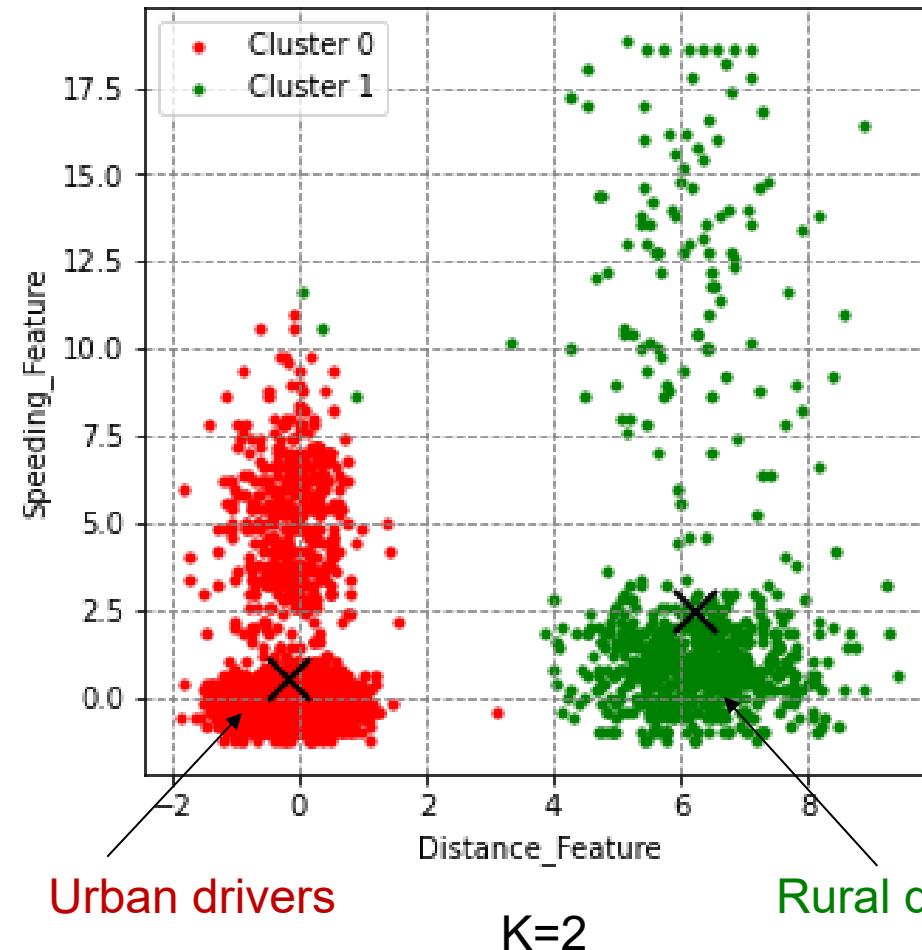
- Assign new data samples to the most related cluster (closest centroid)

```
new_data = ...  
y_pred = kmeans.predict(new_data)
```

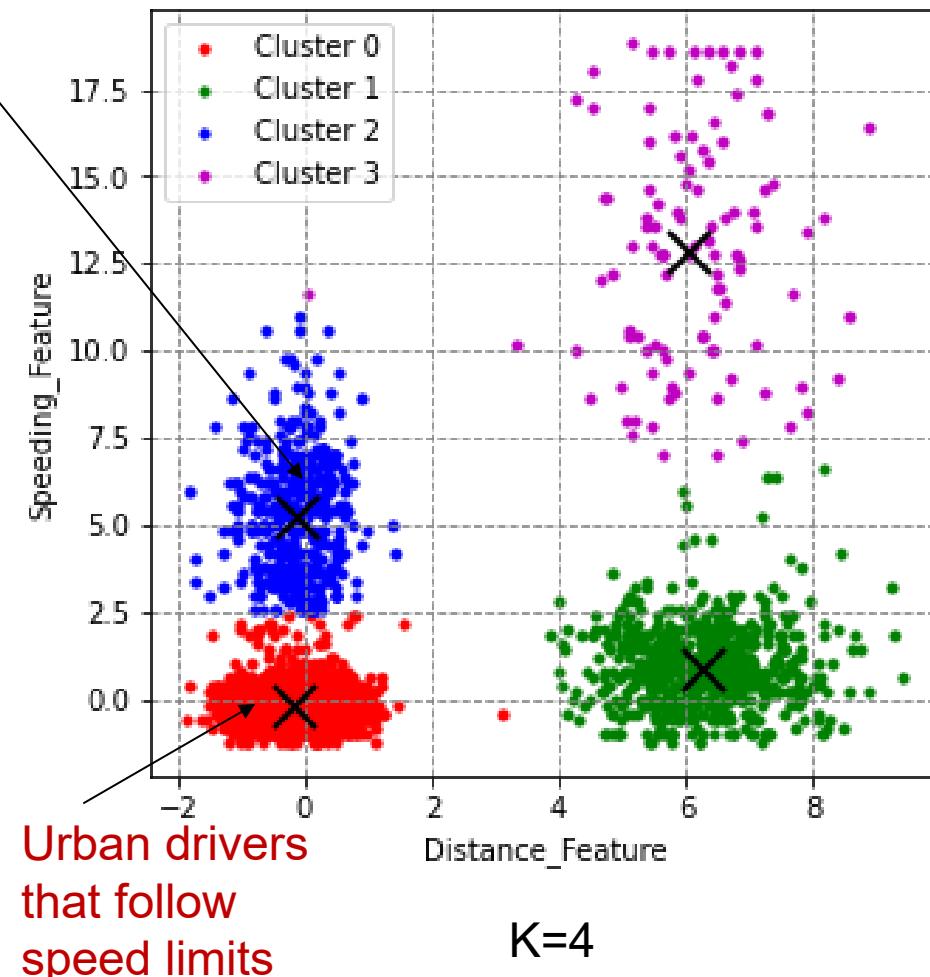


Visualizing results

- Run the K -means clustering algorithm for a range of K values
- Review the results



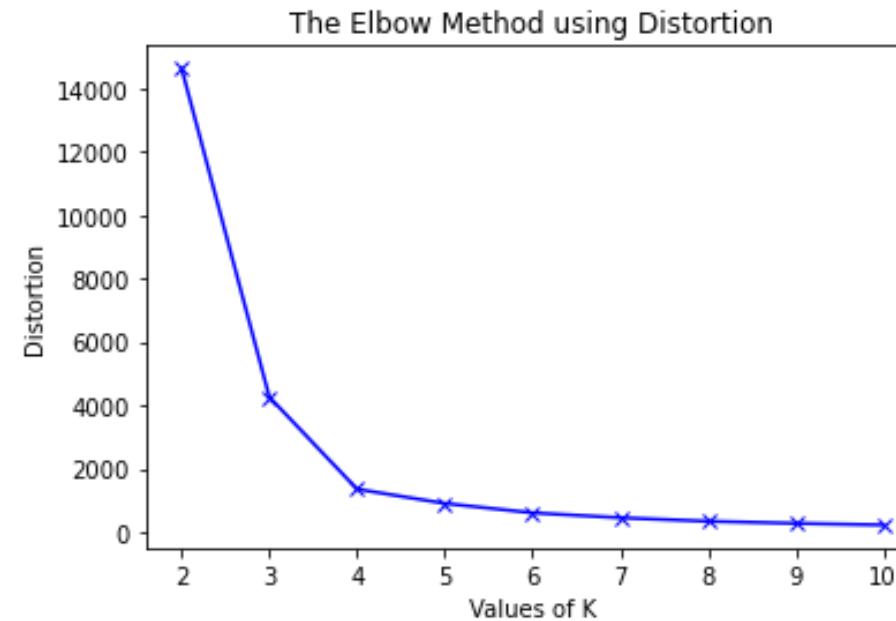
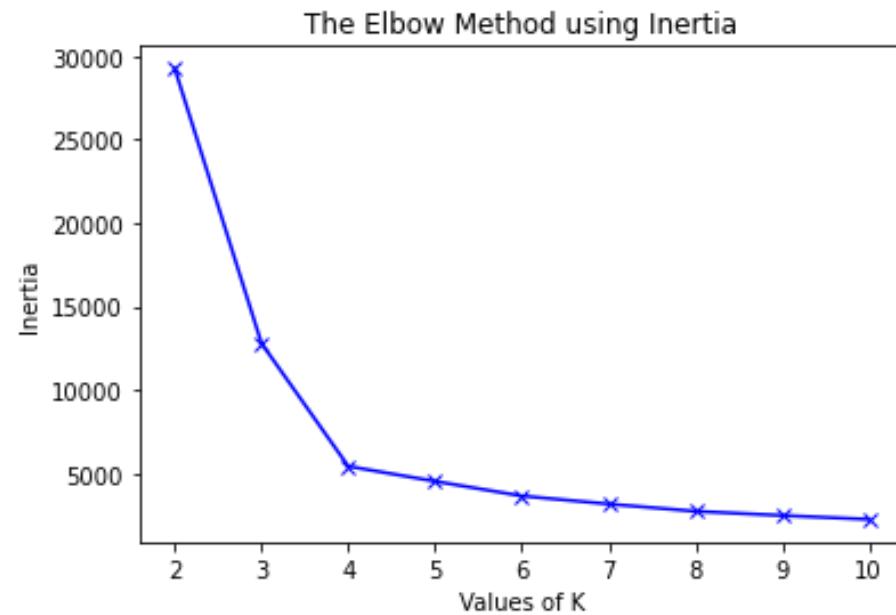
Urban drivers that are speeding frequently





Elbow method

- We run the K-means algorithm for the values of k from 2 to 10 and plot the values of inertia and distortion for each iteration

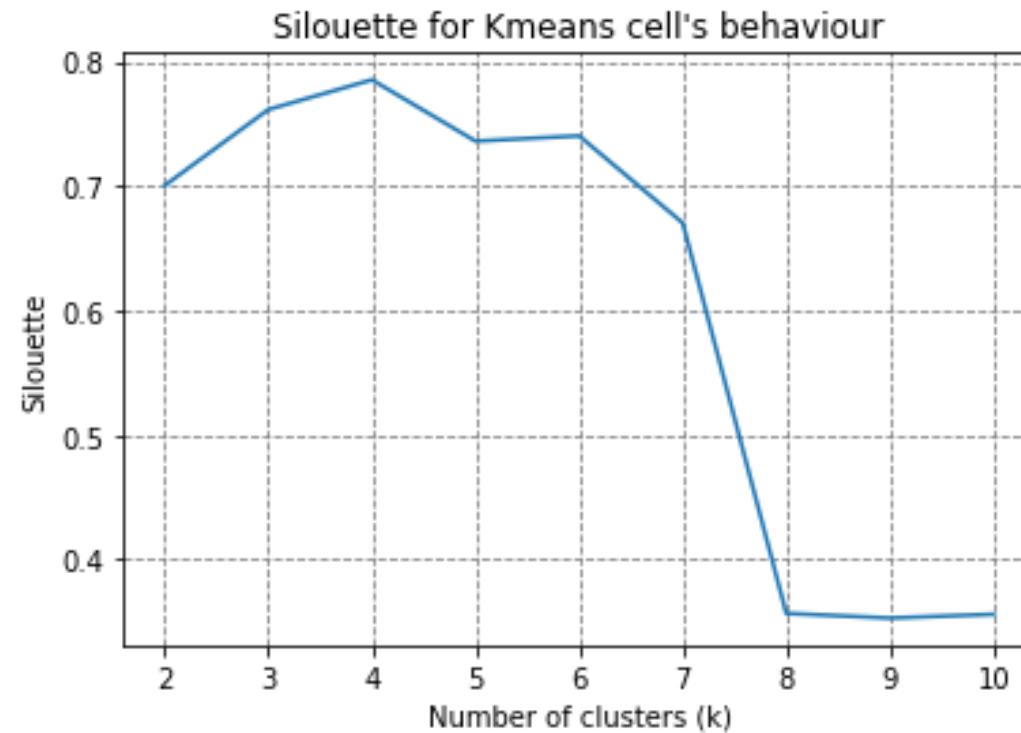


- It seems that the best number of clusters for grouping drivers is 4



Silhouette score

- We run the K-means algorithm for the values of k from 2 to 10 and plot the mean Silhouette score for each iteration



- Silhouette score confirms that the best number of clusters for grouping drivers is 4



Agglomerative clustering

- Perform agglomerative clustering and print labels

```
from sklearn.cluster import AgglomerativeClustering  
# create Agglomerative object and run clustering on the input  
values (X)  
# default k (n_clusters param) → 2  
agglomerative = AgglomerativeClustering (n_clusters=4).fit(X)  
print(agglomerative.labels_) # labels of each sample
```



Clustering for handling missing values

- Clustering (and sometimes classification) can be used to impute missing values in both categorical and numerical features.
- This approach leverages the similarity between data points rather than blindly filling missing values with global statistics (mean, median, or mode).
- **Step 1: Cluster the Data**
 - Temporarily remove the column(s) with missing values that you want to impute.
 - Apply a clustering algorithm (e.g., K-Means) to the remaining data.
 - Use methods like the **Elbow Method** or **Silhouette Score** to determine the optimal number of clusters K .



Clustering for handling missing values

- **Step 2: Assign Missing Values Based on Cluster Membership**
 - Once clusters are formed, each row is assigned to a cluster.
 - For rows with missing values in **categorical columns**, impute using the **most frequent value** of that column among other rows in the same cluster.
 - For rows with missing values in **numerical columns**, impute using the **mean or median** of that column among other rows in the same cluster.
- **Benefits:**
 - Often yields more accurate imputations than simply using global statistics (mean, median, or mode on all values of the column).



Assignment

- Download Lab09 - Clustering Assignment-students-file-with-figures.ipynb from Piazza
- Read comments carefully and implement all requested tasks. You need to replace the None commands with the appropriate command(s).
- Run the file and save results within the same file
- Submit .ipynb file to Moodle by Tuesday 18th of November @ 23.59
 - Login to Moodle: <https://moodle.cs.ucy.ac.cy/course/view.php?id=312> using your UCY credentials
 - Follow “Lab9 Submission” link
 - Upload your .ipynb file



APPENDIX

Linkage methods in Hierarchical clustering



- When we have more than one data point in a cluster, how do we calculate distance between these clusters? To calculate distance we can use any of following linkage methods:
 - Single linkage: considers the minimum distance between two points belonging to different clusters
 - Complete linkage: considers the maximum distance between two points belonging to different clusters
 - Average linkage: considers the average pairwise distance between all points belonging to different clusters
 - Centroid linkage: considers the distance between centroids of clusters
 - Centroid can be calculated as the mean of data points of the cluster
 - E.g. data points: (80, 56), (75, 53), (60, 50), (68, 54) centroid: (70.75, 53.25)



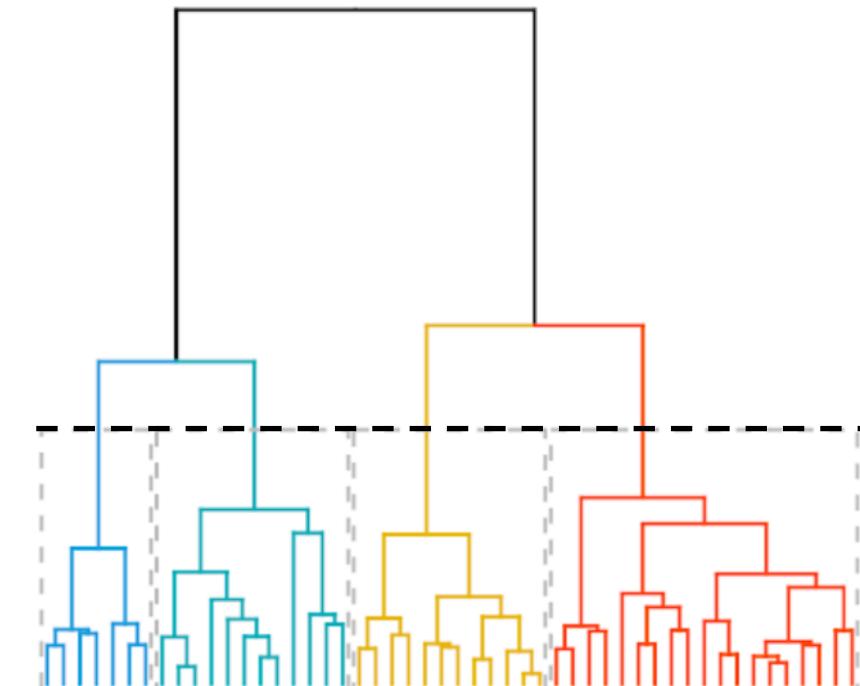
Bisecting K-Means (B-K-Means) vs K-Means

- Bisecting K-Means is a hybrid approach between Divisive Hierarchical Clustering (top down clustering) and K-means Clustering
 - Instead of partitioning the data set into K clusters in each iteration, Bisecting K-Means algorithm splits one cluster into two sub clusters at each bisecting step (by using K-Means) until K clusters are obtained.
- Bisecting K-Means is more efficient when K is large.
 - For the K-Means algorithm, the computation involves every data point of the data set and K centroids. On the other hand, in each Bisecting step of Bisecting K-Means, only the data points of one cluster and two centroids are involved in the computation. Thus, the computation time is reduced.
- Bisecting K-Means produce clusters of similar sizes, while K-Means is known to produce clusters of widely different sizes.

Dendograms for finding best cluster number



- Dendograms **cannot** always tell you how many clusters you should have
- However, if there is an obviously “correct” number of clusters this will often be evident in a dendrogram
- We can plot a dendrogram using [this \(SciPy\) way](#) or [that \(sklearn\) way](#)
 - Since vertical lines represent the distances between merged clusters, we can stop merging clusters when distance among them is long
 - For example, this horizontal cut seems to separate 4 disjoint clusters
- See more details [here](#)



Dendograms for finding best cluster number



- Here is the dendrogram for the drivers dataset which reveals 3 or 4 disjoint clusters

